
Crafter: Towards Automated Reproducible Machine Learning via Agentic Code Generation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Reproducing machine learning research is challenging because many papers do
2 not provide executable code, and key implementation details are often scattered,
3 implicit, or missing. Existing paper-to-code systems improve over direct prompting,
4 but the generated repositories can still miss paper-critical logic or fail at execu-
5 tion time. We propose **Crafter**, an automated paper-to-code pipeline that treats
6 reproduction as a problem of context-calibrated implementation recovery rather
7 than single-pass code generation. Crafter builds an evidence-grounded specifi-
8 cation, resolves missing implementation details before planning, and repairs the
9 generated repository through execution-aware debugging. We evaluate Crafter
10 with PaperBench Code-Dev for implementation faithfulness and a self-designed
11 execution-oriented benchmark for runtime readiness. Across 23 PaperBench papers
12 using the Claude Sonnet 4.6 backend, Crafter achieves an average Code-Dev score
13 of 0.84, improving over Paper2Code by 19.7% and over DeepCode by 22.1%. In
14 execution evaluation, Crafter reaches the contribution-level milestone on all three
15 repositories, compared with two for both baselines, while reducing average repair
16 cycles from 11.0–11.3 to 5.0.

17 1 Introduction

18 Reproducibility is fundamental to machine learning (ML) research, as reproducible implementations
19 enable the community to validate published claims, establish reliable baselines, explore alternative
20 design choices, and build upon prior work. However, empirical evidence shows that a substantial
21 portion of published ML studies are difficult or impossible to replicate due to a lack of transparency,
22 unavailable code and data, inadequate reporting of experimental conditions, and methodological
23 pitfalls such as data leakage [36, 41, 13]. The reproducibility challenge substantially undermines
24 the trustworthiness of ML research and impedes innovation. Recent advances in LLM-enabled
25 code generation (e.g., Codex [7] and Claude [2]) have demonstrated promising coding capabilities
26 across programming languages such as Python, JavaScript, and SQL. Leading technology firms have
27 adopted these tools in daily development [31, 9]. Inspired by LLM-enabled code generation, we
28 ask the research question: *Can we leverage LLMs to generate code that rigorously follows the ML*
29 *literature, reproduces results, designs and executes experiments, and assesses results without human*
30 *intervention?*

31 General-purpose code generation does not directly translate into reliable paper-to-code generation
32 because 1) papers contain underspecified details, e.g., key hyperparameters; 2) Papers describe
33 algorithms at a conceptual or mathematical level; 3) Research implementations often depend on
34 ecosystem-specific behaviors (e.g, distributed training setups) that are not fully captured in the paper
35 text; 4) Papers involve novel abstractions or theoretical constructs that require interpretation before
36 implementation. On the other hand, recent systems such as Paper2Code [37] and DeepCode [24]
37 propose more effective automated multi-stage pipelines for code generation from papers than direct

LLM-based generation. Despite the improved organization of the generated code, two significant gaps remain: the faithfulness to the published algorithms and the executability of the generated code. By examining the limitations observed in Paper2Code and DeepCode outputs, we identify two recurring error types that underlie these gaps. The first is *fabrication-by-extrapolation*, in which the model overextends partial evidence. The second is *fabrication-by-omission*, in which the model fills in missing implementation details with plausible but ungrounded defaults. These errors suggest that paper-to-code generation requires not only a structured pipeline, but also explicit control over what evidence each generation step receives.

We propose **Crafter**, an automated paper-to-code generation pipeline that bridges the gaps in faithfulness and executability through context-calibrated generation. Our key insight is that a faithful reproduction requires finer-grained evidence extraction and explicit recovery of missing implementation information. Instead of treating the paper as a flat document, Crafter organizes it into implementation-oriented evidence. It gathers information from text, figures, tables, and appendices, and identifies important details that are missing or unclear. These gaps are resolved before repository planning, so the code generator receives a more comprehensive implementation plan. The generated repository is then passed through an execution-aware debugging loop that fixes syntax and runtime errors. In this way, Crafter shifts the goal of generation from producing plausible source files to producing verifiable and executable repositories.

We evaluate Crafter with two metrics. First, following prior work, we use the Code-Dev mode of PaperBench [43] to measure code-development faithfulness: whether the generated repository satisfies fine-grained implementation requirements derived from real papers. Second, we build an execution-oriented evaluation benchmark using a three-paper subset of PaperBench to measure how closely the generated repositories resemble real code execution. This evaluation complements PaperBench Code-Dev by testing whether a codebase can move from implementation coverage to executable reproduction. Our results of PaperBench Code-Dev show that Crafter achieves 19.7% and 22.1% higher score than Paper2Code and DeepCode, respectively. For execution evaluation, Crafter reaches the contribution-level milestone for all three papers, compared with two for the two baseline systems.

2 Related Work

Code Generation with Large Language Models. Code generation has become increasingly prominent with the rise of large language models trained on code [18, 15]. Recent developments in this area can be broadly grouped into three main directions: model-level innovations, data-level adaptations, and inference-time refinement strategies. Model-level innovations focus on changing model architecture or the fundamental generation mechanism of LLMs to better understand and generate code. For instance, StarCoder 2 [25] utilizes a Fill-in-the-Middle (FIM) objective, allowing the model to leverage both preceding and succeeding code context to perform precise infilling tasks. DeepSeek-Coder-V2 [58] introduces a Mixture-of-Experts (MoE) [39] architecture and a refined tokenizer to improve efficiency and cross-language understanding while scaling to larger model sizes. CRYSTAL[45] adopts a joint encoder-decoder design that enables unified handling of both natural and programming languages, improving bidirectional reasoning between code and text. Meanwhile, LLaDA [54] replaces the traditional autoregressive Transformer with a diffusion-based generation process, modeling code generation as a denoising task rather than sequential token prediction. On the other hand, data-level adaptations improve models after pre-training by fine-tuning on domain-specific data or using preference optimization. Code Llama [33] and Qwen Coders [52, 17] are fine-tuned on large programming datasets to improve reasoning and correctness, while Magicoder [48] builds open-source instruction data to make generation more controllable and accurate. Finally, inference-time refinement strategies change how LLMs plan, generate, and verify code at runtime. Methods such as AgentCoder [16] and OpenCodeInterpreter [56] use multi-agent collaboration and execution feedback to iteratively test and correct code, while LDB [57] performs step-by-step debugging to catch and fix runtime errors. These methods enhance reliability without retraining the model.

Paper-to-Code Generation. Paper-to-code reproduction is harder than general code generation because papers omit implementation-critical details and depend on ecosystem-specific behaviour rarely captured in the text. Paper2Code [37] and DeepCode [24] address this with multi-stage and

multi-agent pipelines, but both absorb implementation gaps silently into generation and stop at source-level coverage rather than runtime readiness.

3 Methodology

In this section, we introduce the context calibration technique, the Crafter architecture, memory management, and the end-to-end pipeline that includes paper parsing, planning, code generation, and debugging.

3.1 Code Generation with Context Calibration

We formulate the paper-to-code problem as a challenge of *contextual calibration*. We identify two hallucination types in repository-level generation. The first is **fabrication-by-extrapolation**: the model extends partial evidence beyond what the paper supports. For example, if a paper mentions an adapter module, the model may invent an extra routing API or a new hyperparameter that is never specified. The second is **fabrication-by-omission**: the model fills an unspecified implementation choice with its own default. For example, if the paper does not report the learning rate, the model may silently choose one instead of marking it as an unresolved gap.

Crafter addresses these failures through a single principle: every agent invocation operates within a typed, bounded context. When the main risk is extrapolation, Crafter applies **Contraction**, restricting the context to the information needed for the current decision; When the main risk is omission, Crafter applies **Expansion**, adding evidence in a controlled way through a gap-analysis stage that resolves missing details against trusted sources.

Formally, Crafter implements contextual calibration as a spec-centered process. Given a research paper P , the system first builds a structured implementation specification:

$$S = f_{\text{spec}}(P)$$

The specification S serves as the implementation contract, containing the model architecture, loss functions, training procedures, hyperparameters, and other implementation-critical details. To reduce fabrication-by-omission during this stage, f_{spec} applies the Expansion strategy: it identifies information gaps and resolves them against external registries and citation APIs when needed. This helps ground S in paper evidence and trusted external sources rather than LLM intuition.

The code generator then produces an initial codebase:

$$C_0 = f_{\text{code}}(S)$$

To reduce fabrication-by-extrapolation during generation, f_{code} applies the Contraction strategy: the context for each module is bounded to the specifications and dependencies defined in S , preventing the model from inventing symbols or interfaces outside the established contract.

A key design choice is that validation is also derived from the specification, rather than the generated code itself. We construct a set of executable checks:

$$Q = \{q_i\}_{i=1}^m = f_{\text{check}}(S)$$

where each check q_i tests a specific component of the implementation contract, such as unit tests, static analysis, or artifact-level checks. Each check either passes or returns a diagnostic. At repair step t , we collect the set of failing diagnostics:

$$D_t = \{q_i(C_t) : q_i(C_t) \text{ fails}\}$$

If D_t is empty, the codebase is accepted. Otherwise, the diagnostics are used by a repair agent R to iteratively update the code:

$$C_{t+1} = R(C_t, D_t, S)$$

The final output is the first codebase C^* that satisfies all spec-derived checks:

$$q_i(C^*) = \text{pass} \quad \forall i$$

This formulation separates three distinct roles: the **paper** provides evidence; the **specification** defines the implementation contract; and the **checks** convert that contract into executable validation. By decoupling the contract from the implementation, we ensure that C^* must pass all checks produced independently of the code generation process, reducing the probability of producing plausible-looking but non-functional code.




Figure 1: **Crafter architecture.** Ten types of agent across 3 stages pass information through a shared GlobalMemory. The partition agent decomposes the paper into structured partitions, which feed the VLM and gap analysis. Preprocessing outputs are passed to the planning agents, which produce the implementation plan used by the implementation loop.

134 3.2 System Architecture & Memory Model

135 Crafter is organized as a pipeline of multiple specialized agents, as shown in Figure 1 and detailed in
 136 Section 3.3. Unlike prior multi-agent systems that rely on unstructured dialogue [32, 50], Crafter
 137 enforces a fixed inter-agent dataflow mediated by a shared GlobalMemory. The **GlobalMemory** is a
 138 stage-indexed key-value store that serves as the single source of truth across the pipeline, allowing
 139 stages to remain decoupled and re-runnable in isolation. In addition, each agent maintains its own
 140 **LocalMemory**, which records its private trace of reasoning and intermediate attempts. This local trace
 141 is especially important for debugging: the Debug Agent consumes the code generator’s LocalMemory
 142 to recover the intended behavior behind a failed file, helping it distinguish between an incorrect
 143 implementation and an incorrect plan.

144 3.3 Pipeline Design

145 **Structured Partitioning and VLM Analysis.** The process begins by converting the raw PDF
 146 into a typed substrate. A deterministic Partition Agent segments the paper into body, appendix,
 147 references, and preamble. Simultaneously, a Vision-Language Agent performs figure-aware
 148 analysis. Rather than simply describing images, it classifies figures to exclude result plots and
 149 performs structured extraction of architecture diagrams into graph-based records (components, flows,
 150 stages). This phase converts lossy visual information into a typed format that downstream text-based
 151 planners can consume without re-viewing the image.

152 **Gap Analysis with External Resolution.** A dedicated Gap Agent identifies five cate-
 153 gories of implementation-critical omissions: `referenced_architecture`, `baseline_method`,
 154 `dataset_detail`, `missing_hyperparameter`, and `algorithm_gap`. The agent resolves these
 155 gaps through a three-tier cascade. It first checks a local registry for canonical ML datasets, such as
 156 CIFAR [22] and ImageNet [10]. It then performs in-paper lookup by cross-referencing appendices
 157 and footnotes. If the gap remains unresolved, the agent queries the Semantic Scholar [20] API using
 158 title matching, and fills the gap only when the retrieved source passes a string-similarity gate with
 159 OCR-ligature normalization. Gaps that remain unresolved are surfaced as explicit records, making
 160 missing implementation decisions visible rather than silently filled by the model.

161 **Hierarchical Planning and Constraint Validation.** Planning is split between a Main Planning
 162 Agent, which defines the global architecture, and a fan-out of Sub-Planning Agents, which produce

163 module-level specifications. Before code is written, a **constraint validator** enforces orchestrator-
164 module hygiene, reducing extrapolative planning errors by ensuring that entry points such as `main.py`
165 coordinate the workflow without duplicating logic that belongs in implementation modules. A final
166 **cross-module review** pass then checks for interface inconsistencies across parallel-generated sub-
167 plans, such as mismatched return types, and injects reconciliation patches into the subsequent
168 code-generation prompts.

169 **Execution-Aware Code Generation & Debugging** Code generation applies the contraction princi-
170 ple through asymmetric context. **Module-generation** prompts include only the source code of their
171 declared dependencies, which prevents unrelated symbols from leaking across modules. In contrast,
172 **orchestrator-generation** prompts include the full source code of all modules, allowing entry points
173 and training scripts to wire the generated components correctly.

174 The Debug Agent closes the loop by running the codebase in a host subprocess. It utilizes a multi-
175 granularity fix tool (`submit_fix`), which can choose between localized line edits, full-file rewrites,
176 or triggering a dependency-cascade regeneration. If a file fails repeatedly, the system escalates from
177 local repair to re-invoking the planner-conditioned generation for that module and all its importers.

178 4 Experiment Results

179 We evaluate Crafter-generated code repositories using two benchmarks. First, we use the Code-Dev
180 mode of PaperBench [43] to measure code-development faithfulness, i.e., whether the generated
181 repository satisfies the fine-grained requirements of a paper. Second, we measure execution readiness
182 using an execution-oriented subset comprising three papers from PaperBench. These two evaluations
183 separate the implementation coverage from the actual execution. The two evaluations are comple-
184 mentary in that the PaperBench Code-Dev measures implementation coverage, and the execution
185 readiness quantifies practical running capability and repair cost.

186 4.1 Experiment Setup

187 We use PaperBench [43] as the benchmark throughout the research. PaperBench is a benchmark to
188 evaluate AI systems for reproducing ML papers. It consists of 23 papers from ICML 2024. For each
189 paper, PaperBench provides a hierarchical rubric tree, where the leaf nodes describe fine-grained
190 requirements derived from the original paper. Scores are aggregated upward based on assigned
191 weights, with the root score serving as the system’s final score for that paper. In our first evaluation,
192 we focus on the Code Development leaves, which assess whether the generated repository implements
193 the required models, algorithms, datasets, training procedures, baselines, and evaluation logic.

194 We compare Crafter with Paper2Code and DeepCode. All methods take the same input papers and are
195 evaluated with the same PaperBench Code-Dev judge and scoring protocol. For the code-development
196 evaluation, we run all three methods on the 23 papers in PaperBench. For the execution readiness
197 evaluation, we use three papers from PaperBench.

198 We evaluate all methods in two backend settings. In the main setting, each method uses Claude
199 Sonnet 4.6 [3] as its backend, and PaperBench Code-Dev scores are judged with OpenAI o4-mini[30].
200 In the supplementary setting, each method uses MiniMax M2.7[27] as its backend, and scores are
201 judged with OpenAI gpt-4o-mini[29]. The supplementary setting tests whether the effectiveness of
202 Crafter is specific to one backend model or remains consistent across different generators and judges.
203 For execution-oriented evaluation and debugging, we use the Claude Sonnet 4.6 version and run all
204 methods on a NVIDIA GH200 GPU under the same resource budget.

205 We report two groups of metrics. For Code-Dev, we report the average paper score, the median paper
206 score, the leaf criteria passed, and the win count. Each paper score is the root score in its PaperBench
207 rubric tree. The rubric is a weighted hierarchical tree: each Code-Dev leaf criterion is graded by the
208 judge as 0 or 1, and internal nodes aggregate their child scores as weighted sums. Therefore, the final
209 scores are in the range of $[0, 1]$, where 1 indicates that all weighted Code-Dev requirements for that
210 paper are satisfied. The average and median paper scores summarize paper-level implementation
211 quality across the benchmark. The passed leaf criterion measures how many fine-grained rubric
212 requirements are satisfied in papers. The win count measures the number of papers on which a
213 method achieves the highest Code-Dev score among the candidates.

Table 1: PaperBench Code-Dev results on the full 23-paper benchmark. Claude denotes Claude Sonnet 4.6 with o4-mini judging, and MiniMax denotes MiniMax M2.7 with gpt-4o-mini judging. Avg. and Median are paper-level root scores in $[0, 1]$. Passed leaves count the total number of satisfied Code-Dev leaf criteria across all 23 papers. Wins counts the number of papers on which a method obtains the highest score within the same backend setting.

Backend	Method	Avg.	Median	Passed leaves	Wins
Claude	Paper2Code	0.6979	0.7712	2887/3942	4
Claude	DeepCode	0.6844	0.7410	2578/3942	2
Claude	Crafter	0.8354	0.8326	3034/3942	17
MiniMax	Paper2Code	0.4205	0.3837	1440/3942	4
MiniMax	DeepCode	0.4418	0.4406	1233/3942	1
MiniMax	Crafter	0.5875	0.5861	1757/3942	18




Figure 2: **Per-paper Code-Dev: Crafter vs. each baseline.** Each panel pairs one backend (Claude / MiniMax) with one baseline (Paper2Code / DeepCode). A point is one paper at coordinates (baseline, Crafter). Points above the $y = x$ diagonal (blue circles) are papers where Crafter outperforms the baseline; points below (orange / red squares) are papers where the baseline wins. The papers with the largest gain and the largest loss (by $\Delta = \text{Crafter} - \text{baseline}$) are annotated in each panel; corner boxes report wins / 23 and the mean Δ .

For execution-oriented evaluation, we report pre-repair execution milestone, human-repair cycles, source-line edits and the count of repair tiers. These metrics measure how far a generated repository is away from an executable reproduction workflow and how much human repair is required. We define the execution milestones and the repair protocol in Section 4.3.

4.2 Results on PaperBench Code-Dev

Table 1 shows that Crafter consistently outperforms both baselines on the PaperBench Code-Dev. With Claude Sonnet 4.6 as backend, Crafter achieves the highest average and median paper scores, passes the most Code-Dev leaf criteria, and wins on 17 of the 23 papers. This indicates that the improvement is not only reflected in the aggregate score, but also in broad per-paper dominance across the benchmark.

The MiniMax setting confirms the same ranking under a different backend model and judge. Although all methods yield lower absolute scores with MiniMax, Crafter remains the top-performing method across all metrics. This suggests that the relative advantage of Crafter is independent of the chosen backend model.

The two baselines are comparatively close in aggregate performance. Paper2Code passes more leaf criteria in both settings and has more wins than DeepCode in the Claude setting, while DeepCode has a slightly higher average score in the MiniMax setting. In general, the baseline comparison shows that no single baseline dominates the other, while Crafter is consistently ahead of both.

Table 2: Execution protocols. M6 checks liveness: the repository emitted finite, in-range paper-specific metrics. M7 checks the qualitative contribution direction on the same fixture. LCA indicates the paper LCA-on-the-Line

Paper	Fixture	M6 liveness keys	M7 contribution invariant
LBCS	mnist_tiny	acc_lbcs, train_loss	acc_lbcs > acc_uniform
SEMA	cifar_multitask	acc_sema, final_task_loss	acc_sema > acc_baseline
LCA	gaussian_mixture	lca_method, id/ood_top1_error	lca_method < lca_baseline

Per-Paper Analysis. The aggregate results show that Crafter improves the overall performance of Code-Dev. We further analyze the paper-level scores to understand whether this improvement is broad across the benchmark or concentrated in a small number of papers.

Figure 2 illustrates, for each paper, the Code-Dev score of Crafter against each baseline’s score under both backend settings. Points above the $y = x$ line are papers where Crafter wins; points below show baseline wins. In both backend settings, Crafter improves over the baselines on most papers (18–20 out of 23), with mean Δ between +0.14 and +0.17. The losses are on a small set of papers.

4.3 Execution-Oriented Evaluation

PaperBench Code-Dev measures whether a generated repository implements the required paper components in source code. However, source-level coverage does not guarantee that the repository can run as a reproduction artifact. We therefore ask a second question: can the repository launch, complete a small experiment, emit interpretable metrics, and recover the qualitative direction of the paper’s main claim? To answer this question, we build a method-neutral Probe–Plan–Run–Extract harness. We apply the harness to three PaperBench papers whose main empirical claims can be tested with small-scale experiments and compare the execution readiness among the three methods.

4.3.1 Evaluation Design

We evaluate three papers selected from PaperBench: LBCS[51], Self-Expansion[46] (abbreviated SEMA following the authors), and LCA-on-the-Line[40]. We choose papers whose main empirical claims can be tested with small fixture-scale experiments on a single GPU. A fixture refers to a lightweight, paper-specific test setup with reduced data, shortened training, and fixed metrics for checking the core claim. For each paper, we define a fixture-scale protocol with two checks: The M6 liveness check requires the repository to complete the fixture run and emit finite paper-specific metrics. The M7 contribution check requires these metrics to align with the qualitative direction of the major claims in the paper. Table 2 summarizes the protocols. The full paper-specific fixtures, metric keys, and contribution predicates are provided in Appendix D.

We use the same method-neutral Probe–Plan–Run–Extract harness for every paper-method combination. The harness first probes the repository to discover entry points, command-line flags, environment variables, and configuration files. Given this discovered interface, a fixed auxiliary planner chooses commands for a short smoke run and a longer pilot run. The planner is only allowed to use entry points and options that were found in the repository. The harness then runs the repository as a subprocess on the corresponding fixture. After the run finishes, the harness looks for newly created output files and extracts the metrics specified by the paper protocol. The final M6 and M7 predicates are evaluated by the harness, not by the repository or the planner.

We score the three methods using the milestone ladder in Table 3. Milestones M0–M5 measure how far the repository progresses before producing valid paper-specific metrics. M6 indicates that the repository emits valid metrics, and M7 indicates that those metrics satisfy the paper’s contribution-direction predicate. Repositories that do not reach M7 enter a bounded human-repair loop with a 20-cycle cap. We report the pre-repair milestone, the number of repair cycles required to reach M7, source-line edits, and the count of repair tiers. Repair tiers are grouped into S1, S2, and S3, ranging from mechanical surface fixes to localized logic/configuration fixes and architectural or claim-level repairs. Detailed harness rules, error categories, and repair-tier definitions are provided in Appendix D.

Table 3: Execution milestone ladder. Higher milestones indicate deeper progress toward an executable, claim-bearing reproduction.

Milestone	Criterion
M0	No entry point discovered.
M1	Entrypoint found, but did not launch or failed immediately.
M2	Launched, but smoke run did not complete within timeout.
M3	Smoke run completed, but produced no nonempty artifact.
M4	Smoke artifacts exist, but pilot run did not complete.
M5	Pilot run completed, but liveness predicate failed.
M6	Liveness predicate passed, but contribution predicate failed.
M7	Contribution predicate passed.

Table 4: Pre-repair execution outcomes for the first invocation of each generated repository. Each cell reports the milestone reached before any operator repair, with the automatically classified first failure category shown in parentheses. No operator patches are applied in this stage.

Method	LBCS	SEMA	LCA-on-the-Line
Paper2Code	M1 (CLI mismatch)	M5 (CLI mismatch)	M5 (no metrics)
DeepCode	M1 (CLI mismatch)	M1 (missing import)	M1 (CLI mismatch)
Crafter	M5 (syntax error)	M5 (no metrics)	M5 (undefined symbol)

4.3.2 Results.

Execution readiness before repair. Table 4 reports how far each generated repository progresses before any human repair. Without operator intervention, no generated repository reaches the contribution-invariant level (M7). Crafter reaches M5 on all three papers, meaning that each repository launches, completes the pilot run, and produces outputs for the paper-specific liveness check. Paper2Code reaches M5 on two papers, but fails at launch on LBCS. DeepCode fails at or immediately after launch on all three papers. This shows that Crafter is consistently closer to an executable workflow before repair.

Repair distance to M7. Table 5 reports how much human repair is needed to bring each generated repository to M7. Crafter reaches M7 on all three papers with fewer repair cycles and fewer source-line edits than either baseline. More importantly, it requires no S3 architectural-tier patches, indicating that the remaining failures are local rather than structural. The baselines require substantially heavier repair. Paper2Code and DeepCode each reach M7 on two of the three papers, but both require S3 intervention. In practice, these repairs address missing paper-level execution logic, such as absent evaluation tails, unimplemented experimental phases, or emitted metrics that do not correspond to the protocol-defined quantities. These results suggest that Crafter starts closer to the paper-specific contribution checks, whereas the baselines often require architectural repair before the same checks can be evaluated.

The S3 repairs usually reflect missing claim-level execution paths. Baseline repositories often contain paper-named components, but they do not always connect them to the experiments required to test the paper’s major claim. For instance, LBCS baselines miss the comparison between the train-on-coreset and the uniform-baseline, and LCA-on-the-Line baselines miss the synthetic phase used by our executable invariant. We also observe metric mismatches, where a paper-cited function name is present, but the emitted scalar does not match the protocol-defined quantity. These failures are difficult to detect with source-level judging alone, but they dominate the effort needed for executable reproduction.

4.4 Ablation Study

We ablate Crafter on a five-paper subset to isolate the effectiveness of gap filling and debugging. We compare the full Crafter pipeline against two variants: one that keeps gap filling but disables the debug loop, and one that disables both gap filling and debugging. Table 6 reports the per-paper ablation results. The full system achieves the best score on all five papers, with an average score of 0.941. Removing the debug loop reduces the average to 0.871, while removing both gap filling and

Table 5: Aggregate repair distance to M7. “Reached M7” counts the contribution invariant held after repair. “Avg cycles” is the mean number of repair cycles run per cell (cap 20). “ $|\Delta\text{LOC}|$ ” aggregates additions + deletions across all three cells. S1/S2/S3 sum the per-cell distinct patches at each repair tier; “S3 cells” counts cells that required at least one architectural-tier patch.

Method	Reached M7	Avg cycles	$ \Delta\text{LOC} $	S1	S2	S3	S3 cells
Paper2Code	2/3	11.0	322	9	2	3	3/3
DeepCode	2/3	11.3	421	12	6	3	3/3
Crafter	3/3	5.0	232	2	2	0	0/3

Table 6: Per-paper ablation scores on the five-paper subset.

Paper	Crafter	w/o debug	w/o gaps/debug	DeepCode	Paper2Code
BaM[5]	0.980	0.947	0.976	0.946	0.814
PINN[6]	0.951	0.890	0.856	0.869	0.859
All-in-One[12]	0.940	0.866	0.851	0.776	0.867
Mech. Understanding[23]	0.940	0.921	0.896	0.916	0.843
BBOX[44]	0.894	0.731	0.693	0.585	0.722
Average	0.941	0.871	0.854	0.818	0.821

debugging reduces it to 0.854. The ablation result clearly validate the effectiveness of the gap filling and debugging techniques in Crafter.

5 Conclusions

We presented Crafter, a paper-to-code system motivated by the gap between plausible code generation and faithful, executable reproduction. This gap reflects two challenges: papers often omit important implementation choices, and generated repositories may appear complete at the source level while still failing to run or support the paper’s main claims. Crafter makes progress on these challenges by organizing paper evidence into a controlled generation process and by connecting code generation with execution-oriented validation. Across PaperBench Code-Dev and our execution-readiness evaluation, Crafter improves over prior paper-to-code pipelines in implementation faithfulness and requires less repair to satisfy executable contribution checks. Overall, our results suggest that future paper-to-code systems should be designed and assessed on both implementation faithfulness and execution readiness, rather than source-level plausibility alone.

6 Future Works

Although Crafter reduces the repair cycles from 11.0-11.3 to 5.0 in the execution evaluation, there still requires human effort for bug fixing. Further, the scope of Crafter is limited to qualitative evaluation rather than quantitative reproduction. E.g., if a paper claims a 5.5% improvement over baseline, Crafter only confirms the improvement trend, not the exact 5.5% improvement. Another important aspect of reproducible ML research is runtime environment configuration, where the software dependencies is part of the code repository. We will continue work on these three directions with the goal of automated reproducible ML research.

References

- [1] M. S. Albergo, M. Goldstein, N. M. Boffi, R. Ranganath, and E. Vanden-Eijnden. Stochastic interpolants with data-dependent couplings. *arXiv preprint arXiv:2310.03725*, 2023.
- [2] Anthropic. The claude 3 model family: Opus, sonnet, haiku. URL <https://api.semanticscholar.org/CorpusID:268232499>.
- [3] Anthropic. System card: Claude 4.6 sonnet, 2025. URL <https://www-cdn.anthropic.com/78073f739564e986ff3e28522761a7a0b4484f84.pdf>.

- [4] C. Cai, Z. Ye, L. Feng, J. Qi, and F. Liu. Sample-specific masks for visual reprogramming-based prompting. *arXiv preprint arXiv:2406.03150*, 2024.
- [5] D. Cai, C. Modi, L. Pillaud-Vivien, C. C. Margossian, R. M. Gower, D. M. Blei, and L. K. Saul. Batch and match: black-box variational inference with a score-based divergence. *arXiv preprint arXiv:2402.14758*, 2024.
- [6] D. Cai, C. Modi, L. Pillaud-Vivien, C. C. Margossian, R. M. Gower, D. M. Blei, and L. K. Saul. Batch and match: black-box variational inference with a score-based divergence. *arXiv preprint arXiv:2402.14758*, 2024.
- [7] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. de Oliveira Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman, A. Ray, R. Puri, G. Krueger, M. Petrov, H. Khlaaf, G. Sastry, P. Mishkin, B. Chan, S. Gray, N. Ryder, M. Pavlov, A. Power, L. Kaiser, M. Bavarian, C. Winter, P. Tillet, F. P. Such, D. Cummings, M. Plappert, F. Chantzis, E. Barnes, A. Herbert-Voss, W. H. Guss, A. Nichol, A. Paino, N. Tezak, J. Tang, I. Babuschkin, S. Balaji, S. Jain, W. Saunders, C. Hesse, A. N. Carr, J. Leike, J. Achiam, V. Misra, E. Morikawa, A. Radford, M. Knight, M. Brundage, M. Murati, K. Mayer, P. Welinder, B. McGrew, D. Amodei, S. McCandlish, I. Sutskever, and W. Zaremba. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [8] Z. Cheng, X. Wu, J. Yu, S. Yang, G. Wang, and X. Xing. Rice: Breaking through the training bottlenecks of reinforcement learning with explanation. *arXiv preprint arXiv:2405.03064*, 2024.
- [9] K. Z. Cui, M. Demirer, S. Jaffe, L. Musolff, S. Peng, and T. Salz. The effects of generative ai on high-skilled work: Evidence from three field experiments with software developers. *Management Science*, 2026.
- [10] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [11] K. Frans, S. Park, P. Abbeel, and S. Levine. Unsupervised zero-shot reinforcement learning via functional reward encodings. *arXiv preprint arXiv:2402.17135*, 2024.
- [12] M. Gloeckler, M. Deistler, C. Weilbach, F. Wood, and J. H. Macke. All-in-one simulation-based inference. *arXiv preprint arXiv:2404.09636*, 2024.
- [13] O. E. Gundersen, O. Cappelen, M. Mølne, and N. G. Nilsen. The unreasonable effectiveness of open science in ai: A replication study. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 26211–26219, 2025.
- [14] S. Hong, M. Zhuge, J. Chen, X. Zheng, Y. Cheng, J. Wang, C. Zhang, Z. Wang, S. K. S. Yau, Z. Lin, et al. Metagpt: Meta programming for a multi-agent collaborative framework. In *The twelfth international conference on learning representations*, 2023.
- [15] X. Hou, Y. Zhao, Y. Liu, Z. Yang, K. Wang, L. Li, X. Luo, D. Lo, J. Grundy, and H. Wang. Large language models for software engineering: A systematic literature review. *ACM Transactions on Software Engineering and Methodology*, 33(8):1–79, 2024.
- [16] D. Huang, J. M. Zhang, M. Luck, Q. Bu, Y. Qing, and H. Cui. Agentcoder: Multi-agent-based code generation with iterative testing and optimisation, 2024. URL <https://arxiv.org/abs/2312.13010>.
- [17] B. Hui, J. Yang, Z. Cui, J. Yang, D. Liu, L. Zhang, T. Liu, J. Zhang, B. Yu, K. Lu, et al. Qwen2.5-coder technical report. *arXiv preprint arXiv:2409.12186*, 2024.
- [18] J. Jiang, F. Wang, J. Shen, S. Kim, and S. Kim. A survey on large language models for code generation. *ACM Transactions on Software Engineering and Methodology*, 35(2):1–72, 2026.
- [19] X. Jin and X. Ren. What will my model forget? forecasting forgotten examples in language model refinement. *arXiv preprint arXiv:2402.01865*, 2024.

- [20] R. Kinney, C. Anastasiades, R. Authur, I. Beltagy, J. Bragg, A. Buraczynski, I. Cachola, S. Candra, Y. Chandrasekhar, A. Cohan, et al. The semantic scholar open data platform. *arXiv preprint arXiv:2301.10140*, 2023.
- [21] T. Knappe, R. Li, A. Chauhan, K. Chhua, K. Zhu, and S. O’Brien. Semantic self-consistency: Enhancing language model reasoning via semantic weighting. *arXiv preprint arXiv:2410.07839*, 2024.
- [22] A. Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [23] A. Lee, X. Bai, I. Pres, M. Wattenberg, J. K. Kummerfeld, and R. Mihalcea. A mechanistic understanding of alignment algorithms: A case study on dpo and toxicity. *arXiv preprint arXiv:2401.01967*, 2024.
- [24] Z. Li, Z. Li, Z. Guo, X. Ren, and C. Huang. Deepcode: Open agentic coding, 2025. URL <https://arxiv.org/abs/2512.07921>.
- [25] A. Lozhkov, R. Li, L. B. Allal, F. Cassano, J. Lamy-Poirier, N. Tazi, A. Tang, D. Pykhtar, J. Liu, Y. Wei, T. Liu, M. Tian, D. Kocetkov, A. Zucker, Y. Belkada, Z. Wang, Q. Liu, D. Abulkhanov, I. Paul, Z. Li, W.-D. Li, M. Risdal, J. Li, J. Zhu, T. Y. Zhuo, E. Zheltonozhskii, N. O. O. Dade, W. Yu, L. Krauß, N. Jain, Y. Su, X. He, M. Dey, E. Abati, Y. Chai, N. Muennighoff, X. Tang, M. Oblokulov, C. Akiki, M. Marone, C. Mou, M. Mishra, A. Gu, B. Hui, T. Dao, A. Zebaze, O. Dehaene, N. Patry, C. Xu, J. McAuley, H. Hu, T. Scholak, S. Paquet, J. Robinson, C. J. Anderson, N. Chapados, M. Patwary, N. Tajbakhsh, Y. Jernite, C. M. Ferrandis, L. Zhang, S. Hughes, T. Wolf, A. Guha, L. von Werra, and H. de Vries. Starcoder 2 and the stack v2: The next generation. *arXiv preprint arXiv:2402.19173*, 2024.
- [26] M. Malagon, J. Ceberio, and J. A. Lozano. Self-composing policies for scalable continual reinforcement learning. *arXiv preprint arXiv:2506.14811*, 2025.
- [27] MiniMax. MiniMax-M2.7: Built for Agents, coding and beyond, 2025. URL <https://www.minimax.io/news/minimax-m27-en>.
- [28] S. Niu, C. Miao, G. Chen, P. Wu, and P. Zhao. Test-time model adaptation with only forward passes. *arXiv preprint arXiv:2404.01650*, 2024.
- [29] OpenAI. Gpt-4o mini: Advancing cost-efficient intelligence. <https://openai.com/index/gpt-4o-mini-advancing-cost-efficient-intelligence/>, 2024.
- [30] OpenAI. OpenAI o3 and o4-mini system card, 2025. URL <https://cdn.openai.com/pdf/2221c875-02dc-4789-800b-e7758f3722c1/o3-and-o4-mini-system-card.pdf>.
- [31] S. Peng, E. Kalliamvakou, P. Cihon, and M. Demirer. The impact of ai on developer productivity: Evidence from github copilot. *arXiv preprint arXiv:2302.06590*, 2023.
- [32] C. Qian, W. Liu, H. Liu, N. Chen, Y. Dang, J. Li, C. Yang, W. Chen, Y. Su, X. Cong, et al. Chatdev: Communicative agents for software development. In *Proceedings of the 62nd annual meeting of the association for computational linguistics (volume 1: Long papers)*, pages 15174–15186, 2024.
- [33] B. Rozière, J. Gehring, F. Gloeckle, S. Sootla, I. Gat, X. E. Tan, Y. Adi, J. Liu, R. Sauvestre, T. Remez, J. Rapin, A. Kozhevnikov, I. Evtimov, J. Bitton, M. Bhatt, C. C. Ferrer, A. Grattafiori, W. Xiong, A. Défossez, J. Copet, F. Azhar, H. Touvron, L. Martin, N. Usunier, T. Scialom, and G. Synnaeve. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.
- [34] G. Sanchez, H. Fan, A. Spangher, E. Levi, P. S. Ammanamanchi, and S. Biderman. Stay on topic with classifier-free guidance. *arXiv preprint arXiv:2306.17806*, 2023.
- [35] C. Schlarmann, N. D. Singh, F. Croce, and M. Hein. Robust clip: Unsupervised adversarial fine-tuning of vision embeddings for robust large vision-language models. *arXiv preprint arXiv:2402.12336*, 2024.

- [36] H. Semmelrock, T. Ross-Hellauer, S. Kopeinik, D. Theiler, A. Haberl, S. Thalmann, and D. Kowald. Reproducibility in machine-learning-based research: Overview, barriers, and drivers. *AI Magazine*, 46(2):e70002, 2025.
- [37] M. Seo, J. Baek, S. Lee, and S. J. Hwang. Paper2code: Automating code generation from scientific papers in machine learning. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=3DcaUTjdKc>.
- [38] L. Sharrock, J. Simons, S. Liu, and M. Beaumont. Sequential neural score estimation: Likelihood-free inference with conditional score based diffusion models. *arXiv preprint arXiv:2210.04872*, 2022.
- [39] N. Shazeer, A. Mirhoseini, A. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations*, 2017.
- [40] J. Shi, G. Gare, J. Tian, S. Chai, Z. Lin, A. Vasudevan, D. Feng, F. Ferroni, and S. Kong. Lca-on-the-line: Benchmarking out-of-distribution generalization with class taxonomies. *arXiv preprint arXiv:2407.16067*, 2024.
- [41] M. L. Siddiq, A. Islam-Gomes, N. Sekerak, and J. Santos. Large language models for software engineering: A reproducibility crisis. *arXiv preprint arXiv:2512.00651*, 2025.
- [42] J. Singla, A. Agarwal, and D. Pathak. Sapg: Split and aggregate policy gradients. *arXiv preprint arXiv:2407.20230*, 2024.
- [43] G. Starace, O. Jaffe, D. Sherburn, J. Aung, J. S. Chan, L. Maksin, R. Dias, E. Mays, B. Kinsella, W. Thompson, J. Heidecke, A. Glaese, and T. Patwardhan. Paperbench: Evaluating ai’s ability to replicate ai research, 2025. URL <https://arxiv.org/abs/2504.01848>.
- [44] H. Sun, Y. Zhuang, W. Wei, C. Zhang, and B. Dai. Bbox-adapter: Lightweight adapting for black-box large language models. *arXiv preprint arXiv:2402.08219*, 2024.
- [45] T. Tao, J. Li, B. Tan, H. Wang, W. Marshall, B. M. Kanakiya, J. Hestness, N. Vassilieva, Z. Shen, E. P. Xing, and Z. Liu. Crystal: Illuminating llm abilities on language and code, 2024. URL <https://arxiv.org/abs/2411.04156>.
- [46] H. Wang, H. Lu, L. Yao, and D. Gong. Self-expansion of pre-trained models with mixture of adapters for continual learning. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 10087–10098, 2025.
- [47] X. Wang, B. Lin, D. Liu, Y.-C. Chen, and C. Xu. Bridging data gaps in diffusion models with adversarial noise-based transfer learning. In *Forty-first International Conference on Machine Learning*, 2024.
- [48] Y. Wei, Z. Wang, J. Liu, Y. Ding, and L. Zhang. Magicoder: Empowering code generation with oss-instruct. *arXiv preprint arXiv:2312.02120*, 2024.
- [49] M. Wołczyk, B. Cupiał, M. Ostaszewski, M. Bortkiewicz, M. Zając, R. Pascanu, Ł. Kuciński, and P. Miłoś. Fine-tuning reinforcement learning models is secretly a forgetting mitigation problem. *arXiv preprint arXiv:2402.02868*, 2024.
- [50] Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, et al. Autogen: Enabling next-gen llm applications via multi-agent conversations. In *First conference on language modeling*, 2024.
- [51] X. Xia, J. Liu, S. Zhang, Q. Wu, H. Wei, and T. Liu. Refined coreset selection: Towards minimal coreset size under model performance constraints. *arXiv preprint arXiv:2311.08675*, 2023.
- [52] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [53] J. Yang, C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, and O. Press. Swe-agent: Agent-computer interfaces enable automated software engineering. *Advances in Neural Information Processing Systems*, 37:50528–50652, 2024.

- 477 [54] Z. You, S. Nie, X. Zhang, J. Hu, J. Zhou, Z. Lu, J.-R. Wen, and C. Li. Llada-v: Large language
478 diffusion models with visual instruction tuning, 2025. URL [https://arxiv.org/abs/2505.](https://arxiv.org/abs/2505.16933)
479 16933.
- 480 [55] B. Zhao, H. Hajishirzi, and Q. Cao. Apt: Adaptive pruning and tuning pretrained language
481 models for efficient training and inference. *arXiv preprint arXiv:2401.12200*, 2024.
- 482 [56] T. Zheng, G. Zhang, T. Shen, X. Liu, B. Y. Lin, J. Fu, W. Chen, and X. Yue. Opencodeinterpreter:
483 Integrating code generation with execution and refinement, 2025. URL [https://arxiv.org/](https://arxiv.org/abs/2402.14658)
484 [abs/2402.14658](https://arxiv.org/abs/2402.14658).
- 485 [57] L. Zhong, Z. Wang, and J. Shang. Debug like a human: A large language model debugger
486 via verifying runtime execution step-by-step, 2024. URL [https://arxiv.org/abs/2402.](https://arxiv.org/abs/2402.16906)
487 16906.
- 488 [58] Q. Zhu, D. Guo, Z. Shao, D. Yang, P. Wang, R. Xu, Y. Wu, Y. Li, H. Gao, S. Ma, et al.
489 Deepseek-coder-v2: Breaking the barrier of closed-source models in code intelligence. *arXiv*
490 *preprint arXiv:2406.11931*, 2024.

A Prompts

The full system prompts used at every LLM-driven stage of the pipeline are reproduced verbatim from the released source as listing boxes within the relevant methodology subsections of Appendix B. The vision–language stage is covered by Box 2 (classification) and Box B.3.2 (structured extraction); the gap stage by Box 3; planning by Boxes B.3.4–B.3.4 (goal/config extraction, paper analysis, architecture design, architecture replan, sub-planning, and cross-module review); code generation by Box 5 (module files) and Box B.3.5 (orchestrator files); unit-test generation by Box 6; and debugging by Box 7 (tool-using debug loop) and Box B.3.7 (structural rewrite).

B Extended Methodology and Implementation Details

This section provides the full technical description of the Crafter pipeline as detailed in the source methodology documents.

B.1 System Overview

Methodological principle. Crafter is structured around a single discipline applied stage by stage: every LLM call in the pipeline sees a *typed, bounded* context, and the boundary is chosen against the failure mode each call is most prone to. Where the dominant failure mode is fabrication-by-extrapolation—the model inventing a symbol, a class, a hyperparameter, or an import that does not exist—the context is *contracted*: the sub-planner sees only the gaps tagged relevant to its module (§3.3.4), the code-generation prompt sees only the source of files declared as dependencies (§3.3.5), the unit-test prompt sees only an AST-extracted whitelist of importable symbols (§3.3.6), and the debug agent sees the codebase only through bounded `read_file / search_codebase` tool calls rather than as a wholesale dump (§3.3.7). Where the dominant failure mode is the opposite—fabrication-by-omission, in which the LLM produces a plausible value the paper does not state—the context is instead *expanded* with externally retrieved or cross-stage information whose provenance is preserved: the gap-analysis stage surfaces and resolves missing paper-level information against a tiered cascade of trusted sources (§3.3.3), and the planner’s review pass injects cross-module interface findings into every code-generation prompt downstream (§3.3.4). Each of the four contributions below is an instance of one of these two moves, and §3.4 makes the principle’s reach across the pipeline explicit.

Crafter Agent converts a parsed research paper—its sectioned content list and extracted figures—into a runnable PyTorch codebase together with unit tests. The system is organised as a sequence of eight agent types; sub-planning is a single type instantiated once per code module, so the runtime instance count is $8 + N$ for an architecture with N modules. A *partition agent* splits the paper into typed sections (body, appendix, references, preamble) by deterministic parsing of the content list; a *vision–language agent* classifies and structurally extracts architecture- and pipeline-style figures; a *gap agent* identifies information missing from the paper and resolves it through a tiered cascade across three levels of trust—a local dataset registry, an in-paper appendix lookup, and a title-matched Semantic Scholar [20] retrieval—with anything that fails all three retained as an explicit `unresolved_gap` record. A *main planning agent* decomposes the implementation into a file-level architecture and three independent plans (experiments, baselines, dataset loading); a fan-out of *sub-planning agents* produces a per-module specification in parallel; a *code-generation agent* emits each file in dependency order with two passes (modules first, then orchestrator) and a syntax self-correction step; a *unit-test agent* generates tests against an AST-extracted closed symbol table; and a *debug agent* iteratively runs the codebase as a host subprocess and patches errors through three native-tool-call primitives (`read_file`, `search_codebase`, `submit_fix`).

Stages execute sequentially, but planning exploits `asyncio.gather` at three levels—a parallel extraction phase that issues goal-and-config and paper-analysis calls concurrently, three independent planning subtasks (experiments, baselines, dataset loading), and N module sub-planners—to amortise latency across the most expensive part of the pipeline. Inter-agent state lives in a shared `GlobalMemory` keyed by stage and persisted as one file per key with thread-safe `asyncio→background-thread` writes. In parallel, every agent maintains a per-agent `LocalMemory` trace (conversation history, reasoning, intermediate attempts), read by the unit-test and debug agents to recover *why* a particular file was

written a particular way (§3.3.6, §3.3.7) and otherwise persisted as an artefact for post-hoc analysis of agent behaviour.

Relation to prior agentic coding systems. Crafter sits in a small but growing literature on multi-stage and multi-agent code generation. Its closest precedent is PaperCoder [37], which decomposes paper-to-code into a planning–analysis–coding chain. Crafter inherits the staged decomposition and diverges in three substantive ways: (i) a dedicated *pre-planning gap-analysis stage* surfaces paper-level information gaps and resolves them where possible against external sources, rather than absorbing them silently into planning; (ii) planning is *hierarchical and parallel*—three independent planning subtasks fan out into N module sub-planners followed by a cross-module review pass—rather than a single flat planning chain; and (iii) the pipeline includes explicit *debug* and *unit-test* stages that target the gap from “code generated” to “code executes”, which PaperCoder treats as out of scope. Crafter’s typed sequential pipeline also contrasts with the dialogue-driven role-play of ChatDev [32] and AutoGen [50] and with MetaGPT’s SOP-driven waterfall [14]: the methodological commitment in Crafter is *fixed inter-agent dataflow through a shared GlobalMemory* (§3.5), which is what makes individual stages re-runnable in isolation. The debug agent’s tool design is closest to SWE-agent [53] and Aider; Crafter’s `submit_fix` differs in exposing three concurrent granularities of fix in a single tool call and in the planner-conditioned dependency-cascade regeneration triggered after repeated same-file failure (§3.3.7), which lets the agent escalate from local repair to planner-conditioned regeneration without leaving the loop.

The contributions of this work are:

- **Gap analysis with title-matched external resolution.** A dedicated pre-planning stage extracts five categories of information gap (referenced architectures, baseline methods, dataset details, missing hyperparameters, algorithm gaps) and resolves them through a tiered cascade that prefers in-system sources (a curated dataset registry; in-paper appendix lookups) before falling back to a Semantic Scholar query gated by a conjunctive title-similarity check with OCR-ligature normalisation. Gaps that fail every tier are surfaced to downstream agents as explicit `unresolved_gap` records rather than silently filled, so that fabrication is at minimum *visible* in the pipeline state.
- **Hierarchical planning with cross-module review.** Three independent planning subtasks run against a shared architecture, then N module sub-planners run in parallel against gap-filtered context, then a final review pass over the resulting per-module specifications whose output is re-injected into every code-generation prompt downstream. The review pass is intended to surface inter-module interface inconsistencies before code generation begins.
- **Tool-using debug agent with multi-granularity fix primitives.** The debug agent is exposed three native tool calls—`read_file`, `search_codebase`, and `submit_fix`—of which only `submit_fix` is side-effecting. `submit_fix` accepts three concurrent fix modes—localised line-range edits, full-file rewrites, and `files_to_create` that spawn fresh code-generation agents—corresponding to three granularities of repair: edits for localised symptoms, rewrites for whole-file structural problems, and code-generation re-invocation for files the planner missed entirely.
- **Closed-symbol-table unit-test generation.** Test prompts embed an AST-extracted whitelist of every importable top-level symbol from previously generated files, and the test agent is constrained to import only from this whitelist. The whitelist is intended to prevent the failure mode in which LLM-generated tests import symbols the model invented or misnamed.

B.2 Pipeline Orchestration

The pipeline is a hand-rolled `asyncio` program rather than a workflow-engine graph. The top-level `run_pipeline` coroutine awaits each stage in sequence—partition, VLM figure analysis, gap analysis, planning, code generation, debug, and (optionally) unit-test generation—and reports per-stage token counts and accumulated cost between stages. Stage skipping (`-no-vlm`, `-skip-debug`, `-skip-gaps`, `-only-debug`, `-unit-test`) is handled at this level by guarding the corresponding `await`. Within stages, parallelism comes from `asyncio.gather`: the planning stage runs three sequential fan-outs—an extraction phase that issues goal-and-config and paper-analysis calls concurrently, three independent planning subtasks (experiments, baselines, dataset loading), and one

sub-planning agent per code module. No stage runs concurrently with another—the pipeline’s graph is a chain of supersteps, each potentially containing one or more fan-outs.

Agents do not call each other directly. All inter-agent communication flows through a single shared `GlobalMemory` object, instantiated once in `run_pipeline` and threaded into every agent at construction. Each agent reads its inputs from named keys (`partitions`, `figure_analysis`, `architecture`, `module_specs`, `generated_code`, `supplementary_context`, ...) and writes its outputs back to the same store. The store is keyed on disk as one file per key—JSON for structured values, Markdown for free-text values—and is safe across the `asyncio` thread and the background writer threads it spawns: each set serialises its value to a string in the calling task before handing the immutable string to a daemon writer thread that holds a per-store lock. In parallel, every agent maintains a private `LocalMemory` recording its conversation history, reasoning steps, and intermediate attempts. `LocalMemory` uses the same async-safe write pattern as `GlobalMemory`, and is read back in two places: by the unit-test agent, which loads the codegen `LocalMemory` of each file to recover the spec it was written against, and by the debug agent, which can reconstruct the rationale for a particular generated file when diagnosing a failure.

All LLM traffic is mediated by a single client module that exposes three call surfaces—text completion, tool-using completion, and vision—over five interchangeable backends (an Anthropic-compatible `MiniMax` endpoint, Anthropic Claude, OpenAI, DeepSeek, and Gemini), selected by a process-global `set_backend` call. The client handles two failure modes invisibly to its callers. Output continuation: if the model stops with `max_tokens`, the client appends the partial response to the message list and re-issues the request, up to a fixed cap of continuations; if the cap is reached without a stop reason of `end_turn`, the partial output is returned to the caller, who decides whether to accept it or fail the stage. Input truncation: if the request exceeds context, the client repeatedly shrinks the longest user message and retries; the loop terminates either when the call succeeds or when the message reaches a minimum length below which further shrinking would discard required content, at which point the failure is propagated to the caller. Vision calls additionally use exponential backoff on transient connection, timeout, rate-limit, or 5xx errors. A global `TokenUsage` accumulator splits costs by stage and separates VLM tokens from main-model tokens, so per-stage cost can be reported alongside per-stage latency; the per-stage breakdown across backbones is given in §4.

B.3 Stages

B.3.1 Partition

Partition establishes the *typed substrate* on which the rest of the principle operates: every later contraction or expansion is keyed by section role, item kind, or figure index, and those types originate here. The partition agent is the only stage with no LLM call. It consumes a MinerU-produced `*_content_list_v2.json`, in which the paper has already been OCR’d and segmented into a list of pages, each a list of typed items (`title`, `paragraph`, `equation_interline`, `algorithm`, `table`, `image`, `list`, `page_header`, `page_number`). The agent flattens this into a single item stream, drops `page_header` and `page_number` items, and splits the stream into sections at every title item with `text_level == 1`. Each resulting section is tagged with a `section_role` chosen from `{body, appendix, references, preamble}` by regex over the title text (e.g. titles matching “Appendix” or a single-letter “A. Title” pattern become appendix; “References” or “Bibliography” become references). The output written to `GlobalMemory` is two keys: `partitions`, the list of role-tagged sections, and `images`, the figure entries extracted from `image` items so the vision–language stage can iterate over them without re-parsing the content list. Because partitioning is deterministic and idempotent, downstream stages can be re-run against a cached `partitions` value without re-invoking it.

B.3.2 Vision–Language Figure Analysis

Vision–language analysis is the principle’s first instance in the pipeline: a paper’s figure stream is *contracted* by classification (only architecture, pipeline, system-workflow, and algorithm figures survive), then *expanded* into typed structured records—components, flows, stages, inputs, outputs—that downstream stages can splice into a textual prompt without re-showing the image. Most figures in a research paper are not implementation guidance. The vision–language agent’s first job is therefore to filter, not to describe. For each `image` entry surfaced by the partition stage, the agent makes two vision calls. The first is a *classification* call against a prompt that instructs the model to exclude

649 results plots, ablation curves, qualitative examples, and screenshots, and to retain only figures of
650 category architecture, pipeline, system_workflow, or algorithm. The classifier's accuracy
651 is not measured within the pipeline: a misclassified figure either drops from figure_analysis (a
652 true implementation figure rejected) or is fed to the extractor and returns degenerate or contradictory
653 fields (a results plot accepted), in both cases producing low-information output that the planner is
654 robust to under the closed-context principle of §3.1. The second, run only on figures that survive
655 classification, is a *structured extraction* call that returns {category, extracted: {components,
656 flows, stages, inputs, outputs}}—naming each component in the figure, the data-flow edges
657 between them, the sequential stages the figure depicts, and its inputs and outputs. The structured
658 form is what downstream agents consume; the original image is not re-shown to any later stage. The
659 extraction is therefore a typed but lossy compression—typed because the schema is fixed and the
660 planner can rely on field shapes, lossy because the schema cannot represent every detail of the source
661 figure, and unverified within the pipeline because no later stage compares the extracted output back
662 to the source image.

VLM Classification Prompt (src/agents/vlm_agent.py)

```
f"Caption: {caption}\n\n"
"You are a strict figure classifier for a research paper.\n\n"
"Classify this figure into exactly ONE of the following categories:\n"
"- architecture: diagrams showing model layers, blocks, components of a network\n"
"- pipeline: diagrams showing a sequence of processing steps or modules\n"
"- system_workflow: diagrams showing interactions between components, agents, services\n"
"- algorithm: diagrams showing algorithm steps or pseudocode\n"
"- result_plot: experimental results, performance graphs, charts, curves\n"
"- example: input/output examples, generated samples, visualizations\n"
"- other: none of the above\n\n"
"You MUST EXCLUDE (classify as result_plot, example, or other) figures that are primarily:\n"
"- Metrics, curves, score plots, training curves\n"
"- Performance comparisons, ablations, bar charts\n"
"- Qualitative examples, galleries of results, synthesized samples\n"
"- Dataset examples, failure cases, case studies\n"
"- UI screenshots, block catalogs without a full system pipeline\n\n"
"Use ONLY the provided caption text and the visual content of the image.\n"
"Reply with ONLY the category name, nothing else."
```

663

VLM Structured-Extraction Prompt (src/agents/vlm_agent.py)

```
f"Caption: {img.get('caption', '')}\n\n"
"You are an expert vision-language model specialized in understanding figures "
"from computer science and machine learning papers.\n\n"
"Extract structured information from this figure that would help RECONSTRUCT "
"the model, pipeline, or system in code. Be as concrete and implementation-oriented "
"as possible:\n\n"
"1. Components/blocks/modules: list each with a concise name, type, and role\n"
"2. Data flows: source -> target relationships, what data is passed (tensors, features, etc.)\n"
"3. Sequential stages: steps in order, if applicable\n"
"4. Inputs and outputs: data types, shapes, files, models, APIs\n"
"5. Control flow: loops, conditions, branching, interactions\n\n"
"Rules:\n"
"- Prefer information grounded in the VISUAL content; use the caption only as supporting "
"  context\n"
"- Do NOT invent components that are not clearly implied by the figure\n"
"- Be specific about dimensions, layer types, activation functions if visible\n\n"
"Output MUST be valid JSON only. No extra text.\n\n"
"JSON schema:\n"
"{\n"
'  "category": "model_architecture | pipeline | system_workflow | other",\n'
'  "extracted": {\n'
'    "components": [{ "name": "str", "type": "str", "role": "str" }],\n'
'    "flows": [{ "source": "str", "target": "str", "data": "str" }],\n'
'    "stages": [{ "name": "str", "description": "str" }],\n'
'    "inputs": ["str"],\n'
'    "outputs": ["str"]\n'
'  }\n'
"}"
```

664

665 B.3.3 Gap Analysis

666 Gap analysis is the principle’s *expansion* half. Rather than let the LLM fabricate values the paper
 667 does not state, the gap stage surfaces those omissions as explicit, typed records and pulls in additional
 668 information from a tiered cascade of trusted sources, with provenance preserved at every step.
 669 Research papers are written for human readers and routinely omit information a code generator
 670 needs: the precise initialisation of an architecture borrowed from prior work, the exact training
 671 hyperparameters of a baseline cited in a comparison table, the splits and preprocessing of a dataset
 672 assumed to be standard, an algorithmic detail that is “well known” or relegated to an appendix that may
 673 or may not be present. Existing paper-to-code and repository-level code-generation systems share a
 674 common limitation: omissions are absorbed silently into generation rather than surfaced. PaperCoder
 675 [37] folds them into a single planning chain that produces plausible defaults end-to-end with no
 676 explicit record of what was inferred. Generic agentic frameworks (e.g., ChatDev [32], AutoGen [50],
 677 MetaGPT [14]) and tool-using debug agents (e.g., SWE-agent [53], Aider) operate without any paper-
 678 specific notion of an information gap and inherit the same conflation. Neither approach distinguishes
 679 “the paper does not state this” from “the paper states this but we missed it”—a distinction the LLM
 680 cannot recover from after the fact, and which the gap stage is designed to make explicit. The gap
 681 agent is a dedicated stage, run after vision–language analysis and before planning, that surfaces these
 682 omissions explicitly and resolves as many as it can from sources the system trusts. A single LLM call
 683 walks the partitioned paper and emits a list of gaps in five categories—referenced_architecture,
 684 baseline_method, dataset_detail, missing_hyperparameter, and algorithm_gap—with
 685 each gap carrying a structured record {category, description, context_quote, ref_id, ...}.

686 Each gap is then run through a three-tier resolution cascade ordered by cost, with an explicit fallback
 687 for gaps that fail every tier. *Tier 1 (in-system, free)*. A curated local registry of common ML
 688 datasets—including ImageNet, CIFAR-10/100, SQuAD, MMLU, and HumanEval—is checked for
 689 any dataset_detail gap, returning canonical split sizes, class counts, and standard transforms
 690 without an external call. The registry covers the datasets most frequently encountered in the papers
 691 we targeted; gaps for datasets outside the registry fall through to Tier 3. *Tier 2 (in-paper)*. If the
 692 gap carries an internal_ref field, the agent looks up the in-paper appendix the body referenced
 693 but the LLM may not have folded into the original gap context; the partition agent’s section tagging
 694 makes this lookup a constant-time map operation. *Tier 3 (external, network)*. If the gap carries a
 695 ref_id parsed from the bibliography, the agent queries the Semantic Scholar API (trying ArXiv
 696 ID first, then DOI, then a free-text title search) and accepts the result only after a title-similarity
 697 check confirming that the title returned by Semantic Scholar matches the title parsed from the
 698 bibliography under a string-similarity gate combining a SequenceMatcher ratio and a word-overlap
 699 fraction, with OCR ligature corruption normalised before comparison—the normaliser is a small fixed
 700 substitution table that maps MinerU-specific artefacts back to their canonical Unicode form before
 701 similarity is computed (the recurring ± substitution for the ffi ligature is the most common entry;
 702 the table is shipped with the released implementation). The gate is disjunctive—either signal alone
 703 is sufficient—because the two signals catch failure modes that rarely coincide: SequenceMatcher
 704 tolerates word-level substitutions but is sensitive to length mismatch on partial titles, while word-
 705 overlap tolerates character-level OCR corruption but is permissive on short titles, so requiring both
 706 would reject many valid matches. Gaps that fail every tier are surfaced to downstream agents as
 707 explicit unresolved_gap records rather than silently filled, so that fabrication is at minimum *visible*
 708 in the pipeline state.

Gap-Identification System Prompt (src/agents/gap_agent.py)

```
GAP_IDENTIFICATION_SYSTEM = """\
You are an expert ML research analyst. You will receive:
1. The BODY of a research paper (methods, experiments, etc.)
2. The REFERENCES / BIBLIOGRAPHY section

Your job has TWO parts:

## Part 1 -- Parse the references
Extract every reference entry into a structured list. For each, extract:
- ref_id: the citation key as it appears in the paper (e.g. "23", "Hu et al., 2021")
- title: the EXACT title from the bibliography (copy it verbatim)
- authors: author names
- year: publication year
```

709

```

- arxiv_id: arXiv identifier if present (e.g. "2106.09685")

## Part 2 -- Identify implementation gaps
Find every place where the paper references external work that is NEEDED to \
implement the paper's method, experiments, or evaluation, but does not provide \
sufficient detail for a developer to write working code.

Categorize each gap as one of:
1. referenced_architecture: External model/architecture the paper uses or builds on
2. baseline_method: A comparison method the paper evaluates against but does not describe
3. dataset_detail: A dataset where preprocessing/splits/augmentation details are missing
4. missing_hyperparameter: A hyperparameter referenced from another paper
5. algorithm_gap: An algorithm step that references another paper or is garbled/incomplete

For each gap, provide:
- category: one of the five types above
- description: what information is missing (1-2 sentences)
- context_quote: the exact text from the paper where this gap appears
- ref_id: which reference entry this gap refers to (from Part 1). Use null if \
the gap is internal or has no external reference.
- internal_ref: if the paper delegates details to its own appendix or a specific \
internal section (e.g. "see Appendix B for the derivation", "hyperparameters in \
Appendix A.1"), the appendix/section identifier as it appears in the paper \
(e.g. "Appendix B", "A.1", "A"). Use null if the gap does not point to an \
internal section.
- importance: "critical" | "high" | "low"
- what_is_needed: specifically what information would resolve this gap

Do NOT flag standard well-known components (Adam, ReLU, BatchNorm, dropout, etc.).

Output ONLY valid JSON:
{
  "references": [
    {"ref_id": "23", "title": "Exact Title From Bibliography",
     "authors": "...", "year": 2021, "arxiv_id": "2106.09685"}
  ],
  "gaps": [
    {
      "id": "gap_001",
      "category": "baseline_method",
      "description": "...",
      "context_quote": "...",
      "ref_id": "23",
      "internal_ref": null,
      "importance": "critical",
      "what_is_needed": "..."
    }
  ]
}

```

710

711 B.3.4 Hierarchical Planning

712 Planning is where both halves of the principle meet inside one stage: each sub-planner is restricted
713 to the gaps tagged relevant to its module (contraction), while the cross-module review pass injects
714 interface findings into every code-generation prompt downstream (expansion). Planning is also the
715 heaviest stage in the pipeline by token count, and the only stage whose internal control flow is itself
716 parallel. It is split between a *main planning agent* that owns the global view of the codebase and a
717 fan-out of *sub-planning agents*, one per code module, that own the detailed specification of their own
718 files. The main agent runs in five sub-phases.

719 First, an *extraction* phase makes two LLM calls in parallel via `asyncio.gather`: one extracts the
720 project's stated goal and all explicitly given hyperparameters into a `goal_and_config` markdown
721 document, and the other extracts a structured `paper_analysis` covering the equations, algorithms,
722 baselines, ablations, and datasets named in the paper.

723 Second, an *architecture design* phase makes a single LLM call that returns a JSON list of files, each
724 with path, purpose, dependencies, an implements list of equations or algorithms it is respon-
725 sible for, and a boolean orchestrator flag. The output passes through a deterministic constraint
726 validator that forces `orchestrator = true` for any file whose path matches the architecture's
727 own `entry_point` field (defaulting to `main.py` if the LLM omitted it) or the literal basename
728 `reproduce.sh`; any other file's orchestrator flag is taken from the LLM's output, defaulting to

729 `false` if absent. The validator then flags duplicate `implements` claims across files and any orches-
 730 trator that claims `implements` overlap with module files. Conflicts trigger up to a small bounded
 731 number of *replan* calls—a separate LLM call that asks the model to fix the listed conflicts—after
 732 which the validator auto-fixes any residual conflict by stripping `implements` from orchestrator files,
 733 preferring forward progress with a corrected architecture over halting on a hygiene error the validator
 734 can resolve mechanically.

735 Third, the main agent runs three independent planning subtasks in parallel via `asyncio.gather`:
 736 `_plan_experiments` produces an `experiment_plan` that enumerates every experiment the pa-
 737 per runs (proposed method, baselines, and ablations) with the datasets, metrics, and output files
 738 each one writes; `_plan_baselines` produces a `baseline_plan` that gives, per baseline, its file
 739 path, base codebase, training loop, loss formula, key hyperparameters, and differences from the
 740 proposed method, with resolved baseline gaps from the gap agent injected into the prompt; and
 741 `_plan_dataset_loading` produces a `dataset_plan` that names every dataset the paper uses with
 742 its preferred loader (HuggingFace, torchvision, or custom) and any preprocessing the paper specifies.
 743 All three plans are stored under their named keys in `GlobalMemory` and are later threaded into every
 744 sub-planner alongside the architecture and the goal-and-config document.

745 Fourth, the main agent groups the architecture’s files into modules by their top-level directory,
 746 instantiates one sub-planning agent per module, and runs them all concurrently under a second
 747 `asyncio.gather`. Each sub-planner receives the architecture, the goal-and-config document, the
 748 paper analysis, the figure analysis, the three independent plans, and—filtered through the relevance
 749 check described in §3.3.3—only the gaps tagged as relevant to its module. Its output is a markdown
 750 `spec_text` for every file in the module, listing each class with its field types, its `__init__` weight-
 751 initialisation scheme, and per-method step-by-step logic including tensor shapes; every loss formula
 752 written out as an exact equation with variable mappings; and every hyperparameter named with
 753 the exact value from `goal_and_config`. The system prompt instructs the sub-planner to read
 754 hyperparameter values from the configuration document rather than fabricate them, to specify exact
 755 (rather than schematic) formulas, weight-initialisation schemes, and named values from the paper,
 756 and—for any file flagged as an orchestrator—to emit a different spec format consisting only of which
 757 modules to import, the sequence of operations, and the CLI argument parsing, with no algorithm
 758 implementations and no loss formulas. The orchestrator distinction propagates downstream: the
 759 code-generation agent uses it to decide what context to inject (§3.3.5), and the sub-planner is the first
 760 stage to enforce it as a written contract.

761 Fifth, the main agent runs a *review pass*—a single LLM call that reads every module’s `spec_text`
 762 together with the architecture and returns a markdown review flagging interface conflicts: mis-
 763 matched return types, inconsistent function signatures, fields named in one module’s spec but absent
 764 from the file that would have to expose them. The review is stored as `review` in `GlobalMemory`
 765 and re-injected into every code-generation prompt downstream, so file generation begins with the
 766 planner’s reconciliation already in front of the model rather than waiting for inconsistencies to
 767 manifest as runtime errors. The combined effect of the validator on the architecture, the gap-filtered
 768 context per sub-planner, and the review pass is that planning ends with a per-file specification that
 769 has been checked twice—once for orchestrator/module hygiene, once for cross-module interface
 770 consistency—before any code is written.

Goal-and-Config Extraction Prompt (src/agents/planning_agent.py)

```
system = (
    "You are an expert in AI research and reproducing scientific papers.\n\n"
    "Given the sections of a research paper, extract in markdown:\n\n"
    "## Goal\n\n"
    "What is the core method/system? What exactly needs to be implemented?\n\n"
    "Be specific about the model, training procedure, and evaluation.\n\n"
    "## Configuration\n\n"
    "Extract ALL hyperparameters and settings explicitly mentioned in the paper.\n\n"
    "Be EXHAUSTIVE. Include:\n\n"
    "- Training: learning rate, batch size, epochs, optimizer, weight decay, scheduler\n\n"
    "- Model: hidden dims, layers, attention heads, dropout\n\n"
    "- Data: dataset names, preprocessing, augmentation, splits, sample counts\n\n"
    "- Evaluation: metrics, eval frequency, number of samples\n\n"
    "- Loss coefficients: weighting factors, lambda/gamma values\n\n"
    "- Weight init: zero-init, Xavier, Kaiming -- the EXACT scheme\n\n"
```

771

```

    "- Algorithm-specific: step sizes, noise schedules, clipping\n\n"
    "DO NOT FABRICATE -- only use what the paper explicitly provides.\n\n"
    "Output in markdown."
)

```

772

Paper-Analysis Extraction Prompt (src/agents/planning_agent.py)

```

system = (
    "You are an expert in AI research. Given a research paper, extract the following in\n\n"
    "markdown:\n\n"
    "## Equations\n\n"
    "List EVERY numbered equation with its number, formula, and variables.\n\n"
    "## Algorithms\n\n"
    "List EVERY algorithm with its name/number and step-by-step description.\n\n"
    "## Figures and Tables to Reproduce\n\n"
    "List EVERY figure and table showing experimental results, "\n\n"
    "with what data/models are needed to generate each one.\n\n"
    "## Baselines\n\n"
    "List EVERY method the paper compares against with:\n\n"
    "- Name, source (e.g. 'adapts StyleGAN2'), and description.\n\n"
    "## Ablations\n\n"
    "List EVERY ablation study or variant the paper runs.\n\n"
    "## Datasets\n\n"
    "List EVERY dataset used in any experiment with name, split, sample count, and usage.\n\n"
    "Output in markdown. Be thorough -- missing a baseline or dataset means failed reproduction."
)

```

773

Architecture-Design Prompt (src/agents/planning_agent.py)

```

system = (
    "You are an expert software architect specializing in ML/AI research code reproduction.\n\n"
    "Given a research paper's method, goal, configuration, equations, algorithms, and figures to\n\n"
    "reproduce, "\n\n"
    "design the complete code architecture.\n\n"
    "The design must be complete and directly implementable. Avoid unnecessary complexity, "\n\n"
    "and prefer publicly available libraries (PyTorch, numpy, etc.).\n\n"
    "## CRITICAL REQUIREMENTS FOR REPRODUCIBILITY\n\n"
    "The output codebase will be graded on whether it:\n\n"
    "a) Implements code for EVERY equation, algorithm, and method in the paper\n\n"
    "b) Successfully EXECUTES the full training and evaluation pipeline\n\n"
    "c) Produces all figures/tables from the paper's experiments\n\n"
    "Therefore you MUST include:\n\n"
    "- A `reproduce.sh` shell script that runs the ENTIRE pipeline end-to-end "\n\n"
    "(training, evaluation, figure generation) and logs all output to `reproduce.log`.\n\n"
    "- A `main.py` entry point that orchestrates the full pipeline\n\n"
    "- Evaluation/sampling scripts that generate results for EVERY figure and table in the\n\n"
    "paper\n\n"
    "- Each numbered equation in the paper must have a clear corresponding function or code\n\n"
    "block\n\n"
    "- Each algorithm in the paper must have a clear corresponding implementation\n\n"
    "- A `baselines/` module implementing EVERY baseline method the paper compares against. "\n\n"
    "Each baseline gets its own file. If a baseline adapts an existing codebase (e.g. StyleGAN2),\n\n"
    "implement a wrapper that loads/adapts it.\n\n"
    "- All ablation variants of the proposed method must be implementable via config flags or "\n\n"
    "separate scripts -- do not omit ablations\n\n"
    "- A `data/` module that handles downloading and loading EVERY dataset used in the paper\n\n"
    "Output ONLY valid JSON with the following structure:\n\n"
    "```\n\n"
    "{\n\n"
    '  "files": [\n\n"
    '    {\n\n"
    '      "path": "model.py", "purpose": "Main model", "dependencies": ["utils/config.py"], '\n\n"
    '      "implements": ["Eq. 3", "Algorithm 1"], "orchestrator": false},\n\n"
    '      {\n\n"
    '        "path": "main.py", "purpose": "Orchestrate full pipeline", '\n\n"
    '        "dependencies": ["model.py", "training/trainer.py"], "implements": [], "orchestrator":\n\n"
    '          true}\n\n"
    '    ],\n\n"
    '    "modules": ["model", "training", "data", "evaluation", "utils"],\n\n"
    '    "entry_point": "main.py"\n\n"
    '  }\n\n"
    "```\n\n"
)

```

774

```

"Requirements:\n"
"1. Every file must have path, purpose, dependencies, implements, and orchestrator\n"
"   (true/false)\n"
"2. reproduce.sh MUST be included in the files list\n"
"3. Every baseline from the paper must appear as a file under baselines/\n"
"4. Every ablation must be runnable via config flag or separate script\n"
"5. Orchestrator files (entry_point, reproduce.sh, experiment runners) MUST have "\n"
"   "orchestrator: true" and "implements: []". Their job is importing and wiring modules, "\n"
"   "NOT implementing algorithms or equations. All algorithm/equation implementations belong "\n"
"   "in dedicated module files only\n"
"6. No extra text -- output ONLY the JSON"
)

```

775

Architecture-Replan Prompt (src/agents/planning_agent.py)

```

system = (
    "You are a software architect fixing conflicts in a code architecture plan.\n\n"
    "The architecture below has conflicts. Fix them by:\n"
    "1. Removing ALL `implements` entries from orchestrator files (orchestrator=true)\n"
    "2. Ensuring each equation/algorithm is implemented in exactly ONE module file\n"
    "3. Orchestrator files should ONLY import and wire modules together\n\n"
    "Output ONLY the corrected JSON -- same schema, no extra text."
)

```

776

Sub-Planning Agent System Prompt (src/agents/sub_planning_agent.py)

```

system = (
    f"You are an expert researcher and software engineer designing the '{self.module_name}'\n"
    f"module.\n\n"
    "Your task is to produce a COMPLETE and DETAILED implementation specification in markdown "\n"
    "that can be used directly for code generation.\n\n"
    "## Requirements\n"
    "1. STRICTLY align with the paper's methodology\n"
    "2. FOLLOW the overall architecture's interfaces -- do not change the design\n"
    "3. Reference settings from config -- do NOT fabricate values\n"
    "4. For equations: specify the EXACT formula and variable mappings\n"
    "5. For weight init: use the EXACT scheme from the paper\n"
    "6. For hyperparameters: use EXACT values from config\n"
    "7. For loss functions: write the full formula\n\n"
    "## Output Format (Markdown)\n"
    "For each file in this module, write:\n\n"
    "### `file_path.py`\n"
    "***Imports**": list imports\n"
    "***Classes**":\n"
    "- `ClassName(base_class)`: description\n"
    "- `__init__(args)`: what to initialize, weight init scheme\n"
    "- `method(args) -> return_type`: step-by-step logic, tensor shapes, equations\n"
    "***Functions**":\n"
    "- `func(args) -> return_type`: description\n\n"
    "Include tensor shapes (e.g. `[B, C, H, W]`), exact formulas, and default values.\n"
    "Be detailed enough that a developer can implement without reading the paper.\n\n"
    "## Module-specific rules\n"
    "If this module is 'baselines': spec EVERY baseline the paper compares against.\n"
    "If this module is 'data': spec download for EVERY dataset, prefer HuggingFace/torchvision.\n"
    "If a file is marked as `orchestrator: true` in the architecture, its spec should ONLY\n"
    "cover:\n"
    "- Which modules to import and from where\n"
    "- The sequence of operations (data loading -> training -> evaluation -> output)\n"
    "- Command-line argument parsing if applicable\n"
    "- It should NOT specify algorithm implementations, loss formulas, or model architectures -- "\n"
    "those belong in the module files. The orchestrator only imports and calls them."
)

```

777

Cross-Module Review-Pass Prompt (src/agents/planning_agent.py)

```

system = (

```

778

```

    "You are a senior code reviewer performing a final consistency check before
    implementation.\n\n"
    "Given module specifications from multiple independent planning agents, verify that all "
    "interfaces are consistent and compatible. Check for:\n\n"
    "1. Signature mismatches between modules\n"
    "2. Return type conflicts\n"
    "3. Missing dependencies\n"
    "4. Inconsistent data structures\n"
    "5. Circular dependencies\n"
    "6. Config conflicts\n\n"
    "For each conflict found, provide a specific resolution.\n"
    "Output in markdown."
)

```

779

780 B.3.5 Code Generation

781 Code generation applies the contraction principle file by file: each module’s prompt sees only the
 782 source of files it explicitly declared as dependencies, and orchestrators alone are given the full
 783 module source—a deliberate asymmetry between the modules’ “declared neighbourhood” and the
 784 orchestrators’ “name-the-symbols-across-the-codebase” responsibility. Code generation is one agent
 785 invocation per file, run in two passes ordered by the dependency graph the planner produced. The
 786 first pass walks every non-orchestrator file in dependency order—utilities and primitives before
 787 the modules that depend on them—and emits one Python file per call. The second pass walks
 788 the orchestrator files (the entry point and `reproduce.sh`) once every module is on disk, so that
 789 orchestrators can be conditioned on the *actual* generated source of every module they will eventually
 790 call into rather than on its planned spec. The second pass is necessary because orchestrators reference
 791 the *actual* symbols of every module they call into; conditioning them on the planner’s spec rather
 792 than on the realised source produces symbol-name mismatches whenever a module’s planned spec
 793 drifts from its realised source—a failure mode the two-pass design was introduced to remove. The
 794 context assembled for each invocation is deliberately asymmetric. Every file’s prompt receives
 795 the architecture, the file’s own `module_spec`, the goal-and-config document, the paper analysis,
 796 the planner’s review, and the gap entries relevant to the file as identified by the relevance check
 797 in §3.3.3. On top of this, a *module* file receives only the previously generated source of the files
 798 it explicitly declared as dependencies, while an *orchestrator* file receives the full source of every
 799 module—because an orchestrator’s correctness depends on naming the right symbols across the entire
 800 codebase, while a module’s correctness depends only on its declared neighbourhood.

801 The system prompt requires complete code (no TODO, no pass placeholders, no stubs), type hints on
 802 every argument and return, identifiers that match the symbols in previously generated files exactly,
 803 and an explicit injunction that hyperparameters be read from configuration rather than fabricated.
 804 Once the LLM returns, the agent attempts to compile the file with `ast.parse`; if compilation raises
 805 a `SyntaxError`, it makes one self-correction call passing the original output and the parser’s error
 806 message back to the model and uses whichever of the two responses parses, falling back to the raw
 807 output if neither does. `reproduce.sh` receives deterministic post-processing rather than further
 808 LLM intervention: it is rewritten to guarantee a bash shebang, set `-e`, an `exec >>(tee -a`
 809 `reproduce.log) 2>&1` redirect, and a `python -m pip install` step in place of bare `pip`, so
 810 that downstream debug runs always execute under the same logged, fail-fast shell. Each successfully
 811 written file is committed both to disk under the codebase output directory and to `generated_code`
 812 in `GlobalMemory`, where it is immediately visible to the next invocation in dependency order and,
 813 later, to the unit-test agent and the debug agent.

Code-Generation Agent System Prompt – module files (`src/agents/codegen_agent.py`)

```

return (
    "You are an expert researcher and software engineer with a deep understanding of "
    "experimental design and reproducibility in scientific research.\n\n"
    "You will receive a file specification, the overall architecture, configuration, and any "
    "previously generated code files. Your task is to write code that reproduces the paper's "
    "method.\n\n"
    "The code you write must be elegant, modular, and maintainable, adhering to Google-style "
    "guidelines.\n"
)

```

814

```

"The code must strictly align with the paper's methodology, experimental setup, and
evaluation metrics.\n\n"
f"You MUST write only the file: {self.file_path}\n\n"
"## Rules (FOLLOW ALL OF THESE)\n\n"
"1. **Only ONE file**: Do your best to implement THIS ONLY ONE FILE\n"
"2. **COMPLETE CODE**: Your code will be part of the entire project, so implement complete, "
"reliable, reusable code. No TODOs, no placeholders, no '...', no 'pass' in non-abstract
methods\n"
"3. **Set default value**: If there is any setting, ALWAYS SET A DEFAULT VALUE, "
"ALWAYS USE STRONG TYPE AND EXPLICIT VARIABLE\n"
"4. **Follow design**: YOU MUST FOLLOW the data structures and interfaces from the spec. "
"DONT CHANGE ANY DESIGN. Do not use public member functions that do not exist in the spec\n"
"5. **CAREFULLY CHECK** that you don't miss any necessary class/function in this file\n"
"6. **Import first**: Before using an external variable/module, make sure you import it
first. "
"Avoid circular imports\n"
"7. **Write out EVERY CODE DETAIL**: Don't leave any TODO or unimplemented section\n"
"8. **REFER TO CONFIGURATION**: You must use configuration from the provided config. "
"DO NOT FABRICATE any configuration values -- only use what is explicitly provided\n"
"9. **Type hints**: Include type hints for all function arguments and return values\n"
"10. **Previously generated files**: Reference the already generated files for correct
imports "
"and interface compatibility\n"
"11. **EXACT paper values**: Weight initialization, hyperparameters, and loss coefficients "
"MUST match the paper EXACTLY. If the spec says zero-init, use `nn.init.zeros()` -- "
"do NOT substitute Xavier or Kaiming. If gamma=5, hardcode 5 as the default, not 1.0\n"
"12. **Equation fidelity**: If this file implements a paper equation, the code MUST compute "
"every term in that equation. Do not simplify or omit terms\n"
"13. **DO NOT DUPLICATE**: If a class or function already exists in a previously generated
file, "
"IMPORT it -- do NOT redefine it\n\n"
"## Output Format\n"
"Write the code with triple quotes. Output format:\n"
f"```python\n"
f"## {self.file_path}\n"
f"...\n"
f"```"
)

```

815

Code-Generation Agent System Prompt – orchestrator files (src/agents/codegen_agent.py)

```

return (
    "You are writing an ORCHESTRATOR file. Your ONLY job is to import from existing modules "
    "and wire them together.\n\n"
    f"You MUST write only the file: {self.file_path}\n\n"
    "## CRITICAL RULES\n\n"
    "1. **IMPORT, DON'T IMPLEMENT**: You MUST NOT define any model, trainer, dataset, loss, "
    "evaluation, or algorithm logic. Every class and function you need already exists in the "
    "generated module files provided below. IMPORT them.\n\n"
    "2. **Verify interfaces**: For every function call you make, verify the import path and "
    "function signature against the actual source code provided. Do not guess at interfaces.\n\n"
    "3. **No duplication**: If you find yourself writing >20 lines of logic that isn't "
    "argument parsing or control flow, STOP -- you're reimplementing something that belongs "
    "in a module file. Import it instead.\n\n"
    "4. **COMPLETE CODE**: No TODOs, no placeholders, no 'pass' in non-abstract methods\n\n"
    "5. **Set default values**: ALWAYS SET DEFAULT VALUES for settings, "
    "ALWAYS USE STRONG TYPE AND EXPLICIT VARIABLE\n\n"
    "6. **Follow design**: Follow the architecture spec's interfaces exactly\n\n"
    "7. **REFER TO CONFIGURATION**: Use config from the provided config. "
    "DO NOT FABRICATE any configuration values\n\n"
    "8. **Type hints**: Include type hints for all function arguments and return values\n\n"
    "9. **Previously generated files**: Reference the provided source code for correct imports. "
    "Use the EXACT class names, function names, and argument signatures from the source.\n\n"
    + reproduce_rules +
    "\n## Output Format\n"
    "Write the code with triple quotes. Output format:\n"
    f"```python\n"
    f"## {self.file_path}\n"
    f"...\n"
    f"```"
)

# `reproduce_rules` is appended above only when self.file_path == "reproduce.sh":
reproduce_rules = (

```

816

```

"\n## reproduce.sh rules\n"
"- The script MUST start with EXACTLY these lines (no custom logging, no argument parsing):\n"
" ```\n"
" #!/bin/bash\n"
" set -e\n"
" exec > >(tee -a reproduce.log) 2>&1\n"
" python -m pip install -r requirements.txt\n"
" ```\n"
"- ALL output MUST go to `reproduce.log` via the exec redirect above\n"
"- Do NOT create custom log files, log functions, or argument parsers\n"
"- Do NOT use bare `pip` -- always use `python -m pip`\n"
"- Run EVERY experiment: every model variant, every dataset, every baseline\n"
"- Loop over ALL source->target pairs, not just one default\n"
"- Save generated images to `results/images/`, metrics to `results/metrics.json`\n"
"- The script MUST write files to disk -- a no-op script scores 0\n"
"- Must be runnable with `bash reproduce.sh` from the codebase dir, no arguments\n"
"- Print progress: `echo '=== Running experiment X ==='\n"
"- Keep it SIMPLE -- a flat sequence of commands, not functions/stages/parsers\n"
)

```

817

818 B.3.6 Unit-Test Generation

Unit-Test Agent System Prompt (src/agents/unittest_agent.py)

```

system = (
    "You are an expert Python test engineer. Your task is to analyze a source file, "
    "break it into independently testable code blocks (functions, methods, classes), "
    "and write thorough unit tests for each block.\n\n"
    "## Process\n"
    "1. Read the source file and identify every testable unit: functions, class methods, "
    "   property logic, data transformations, etc.\n"
    "2. For each testable unit, write one or more test cases covering:\n"
    "   - Normal/happy path\n"
    "   - Edge cases (empty inputs, boundary values, None, etc.)\n"
    "   - Error conditions (if applicable)\n"
    "3. Use `unittest` (standard library). Use `unittest.mock` for external dependencies "
    "   (file I/O, network, GPU operations, heavy ML training loops).\n\n"
    "## Rules\n"
    "- Mock heavy dependencies (torch training loops, file downloads, API calls) but test "
    "  the actual logic (data processing, shape transformations, config parsing, etc.)\n"
    "- For ML code: test tensor shapes, dtype handling, forward pass output shapes, "
    "  loss computation shapes -- use small random tensors, NOT real data\n"
    "- Every test must be runnable without GPU, without real data, without network access\n"
    "- Import ONLY symbols listed in the 'Exact importable symbols' section. "
    "- If a name you'd want isn't there, work around it -- do NOT invent class or function "
    "  names.\n"
    "- Write COMPLETE test code -- no TODOs, no placeholders\n"
    "- Each test method should test ONE thing and have a descriptive name\n"
    "- Add a brief comment above each test class explaining what code block it tests\n\n"
    "## Output Format\n"
    "```\npython\nimport unittest\n...\n```\n\n"
    "Output ONLY the test file code, nothing else."
)

```

819 Unit-test generation is the strictest application of the contraction principle in the pipeline: the test
820 prompt sees only an AST-extracted whitelist of importable symbols from previously generated files,
821 with everything else excluded by construction. The stage is opt-in, enabled with `-unit-test`, and
822 runs after code generation and before debug. It produces one `tests/test_<module>.py` per non-
823 test, non-`__init__` Python file in the codebase. For each target file the agent does two things that are
824 not done elsewhere in the pipeline: it loads the *codegen LocalMemory* of that file—the conversation,
825 reasoning trace, and intermediate attempts of the agent that originally wrote the file—and it builds a
826 closed table of exactly which symbols the test is allowed to import. The *LocalMemory* load means
827 each test is conditioned on the same spec the file was written against rather than only on the file's
828 final source, so the test reflects the planner's intent (a forward pass that produces a particular shape,
829 a loss in a particular form) and not just whatever the codegen agent happened to compile to. The

unit-test agent and the debug agent (§3.3.7) are the only two cross-agent consumers of LocalMemory in the system, and both read codegen LocalMemory rather than re-deriving spec from source.

The generated tests are produced as a code-generation artefact rather than as a within-pipeline correctness signal: the debug stage (§3.3.7) runs only the entry point declared in the architecture, not the test suite. We surface this distinction explicitly because it bears on what the unit-test stage’s contribution is—a structured, importable artefact that downstream consumers (e.g. a CI pipeline or a human reproducer) can run, *not* a hidden validator the pipeline relies on for its own success criterion.

The closed symbol table is a fabrication-by-extrapolation filter. The agent walks the AST of every previously generated file and extracts every top-level class and function name, excluding underscore-prefixed and `test_*` symbols, mapping each module path to the set of names it exports. This whitelist is included in the test-generation prompt as a section labelled “Exact importable symbols”, with an explicit instruction that the test may import only from that section. The whitelist is intended to prevent the failure mode in which LLM-generated tests import symbols the model invented or misnamed—a class that does not exist, a helper that was never written, a renamed-but-not-renamed function. The system prompt additionally requires the use of `unittest` and `unittest.mock` from the standard library, mandates that GPU operations, file I/O, and network calls be mocked, and instructs the model to test tensor shapes, dtypes, and loss output shapes against small random inputs rather than against trained weights or real datasets. Every test must be runnable on a CPU-only machine without internet access. The generated tests are written to disk under `codebase/tests/` and recorded in `generated_tests` in GlobalMemory; if all expected test paths already exist on disk, the stage skips itself, making it cheap to re-invoke after a partial run.

B.3.7 Debug

Debug applies the contraction principle through *on-demand context*: rather than receiving the full codebase up front, the agent is given two read-only tools (`read_file`, `search_codebase`) and must request the lines it needs to see as the diagnosis unfolds. The debug stage closes the pipeline by running the generated codebase and iteratively patching whatever the run uncovers. Each iteration begins by spawning a host Python subprocess against the entry point declared in the architecture, with the codebase directory as its working directory and a 120-second wall-clock timeout. An exit code of zero terminates the loop with success. A 120 s wall-clock timeout terminates the loop with success only if the captured subprocess output contains at least one liveness token—a case-insensitive regex match for one of `epoch`, `step`, `iter(ation)`, `batch`, `loss`, or `train(ing)` adjacent to a digit—confirming that execution reached the training loop before the timer fired. Timeouts without a liveness token are treated as failures and re-enter the triage path described below. We chose this criterion to separate two outcomes the bare exit code cannot distinguish: a codebase that loaded its model, built its dataloader, and entered training (which we count as a successful run, since the goal of the stage is reaching real work, not waiting out an epoch) from one that hangs in import, deadlocks during dataset construction, or stalls inside an empty loop (which we do not). The fraction of debug runs that succeed via clean exit, via timeout-with-liveness, and via no-success is reported per backbone in §4. A non-zero exit drops into a triage step that parses the captured traceback with regular expressions to identify the *crash file* (the last `.py` referenced in the traceback frames), the line number, and the full set of files mentioned anywhere in the trace. The loop runs up to `max_iterations` rounds (default 20, configurable via `-debug-iterations`), bounded to keep token co-utilisation within limits.

Before invoking the LLM, the agent attempts two cheap recoveries. If the traceback matches `No module named 'X'`, the agent shells out a `pip install X` and re-runs the codebase without consuming an LLM call—most ML papers reach for the same dozen libraries and the host environment cannot enumerate them in advance. If the same error message has now occurred on the same file across repeated iterations, the agent assumes that the LLM keeps patching the wrong file and breaks the loop by triggering a *dependency cascade regeneration*: it scans `generated_code` for files that import the crashed module, re-invokes the code-generation agent for each importer, and finally re-invokes it for the crash file itself, replacing localised LLM edits with a fresh planner-conditioned generation. We trigger after repeated failure on the same file because the LLM-driven path, once converged on a wrong patch, tends to converge on it again on the next iteration—local edits cannot fix structural problems, and burning further iterations on the same file is dominated in expectation by re-invoking the planner-aware code-generation path on its dependents. Both fast paths exist

885 because the LLM-driven path that follows is the most expensive operation in the system, and most
886 repeated-failure traces in practice indicate a structural problem that local edits cannot fix.

887 If neither fast path applies, the agent invokes the main LLM with native tool calling.
888 The model is exposed three tools. `read_file(path)` returns numbered lines from the
889 codebase. `search_codebase(query)` returns grep-style hits with paths and line numbers.
890 `submit_fix(analysis, edits?, file_rewrites?, files_to_create?)` is the only side-
891 effecting tool, must be called exactly once, and accepts three mutually compatible kinds of fix: an
892 edits array of {path, start_line, end_line, new_code} records intended for localised line-
893 range changes, a `file_rewrites` array of {path, content} records that replace whole files when
894 an edit would be more disruptive than a rewrite, and a `files_to_create` array of paths that spawn
895 fresh code-generation agents for files the planner missed entirely. Edits are deliberately the cheapest
896 mode because most fixes the LLM proposes are local (a missing import, an off-by-one, a wrong call
897 signature); rewrites and `files_to_create` exist for the structural cases edits cannot reach. Edits are
898 applied line-range by line-range with re-indentation and a per-edit `ast.parse` syntax check; rewrites
899 overwrite whole files; new files are dispatched to code generation under the same dependency-aware
900 context-assembly described in §3.3.5. Every fix is appended to a `fix_history` recorded in the
901 debug agent's LocalMemory, so a post-hoc reader can reconstruct what the system tried, in what
902 order, and which fix finally let the codebase run. The stage returns when an iteration ends in success,
903 or when `max_iterations` is reached without one—at which point the failure is surfaced with the
904 last traceback rather than silently dropped.

Debug Agent System Prompt (src/agents/debug_agent.py)

```
system = (  
    "You are an expert Python/ML debugger. You have been given an error traceback "  
    "and all codebase files involved. Your job is to find and fix the bug.\n\n"  
    "## How to work\n"  
    "1. Read the FULL traceback -- the bug may not be in the last file. Trace the call chain "  
    "to find the ROOT CAUSE. For example, if main.py calls download.py and download.py crashes, "  
    "fix download.py, not main.py\n"  
    "2. If you need to check another file, use read_file or search_codebase\n"  
    "3. Once you understand the bug, call submit_fix with your fix\n\n"  
    "## submit_fix options\n"  
    "- `edits`: Small targeted line-range changes (under ~15 lines)\n"  
    "- `file_rewrites`: Complete file content -- use for structural bugs, method rewrites, "  
    "or when edits have failed before. Provide the FULL file content.\n\n"  
    "## Rules\n"  
    "- **Fix the RIGHT file** -- the file where the bug originates, not just the file that  
    crashed. "  
    "If the error is a bad argument passed from caller.py to callee.py, fix the caller\n"  
    "- You have at most 5 tool calls. Use them wisely -- don't read files you don't need\n"  
    "- You MUST call submit_fix exactly once\n"  
    "- Only fix codebase files -- NEVER try to fix torch, numpy, or other system packages\n"  
    "- For shape mismatches: check spatial dimensions at each step\n"  
    "- For shape mismatches with nn.Linear on [B,C,H,W] tensors: apply channel-wise "  
    "(permute to [B,H,W,C], reshape to [B*H*W,C], apply Linear, reshape back). "  
    "NEVER flatten to [B,C*H*W]\n"  
    "- If previous fixes failed, try a fundamentally different approach\n"  
    "- Line numbers in edits refer to the ORIGINAL file\n"  
    + rewrite_guidance  
)
```

905

Debug Agent Structural-Rewrite Prompt (src/agents/debug_agent.py)

```
system = (  
    "You are an expert Python/ML engineer. A generated file has a structural bug that "  
    "could not be fixed with small edits. You must REWRITE THE ENTIRE FILE from scratch.\n\n"  
    "## Rules\n"  
    "- Write COMPLETE, WORKING code -- no TODOs, no placeholders\n"  
    "- The file must be compatible with the dependent files shown below\n"  
    "- Match all function signatures and class interfaces exactly\n"  
    "- Fix the bug described in the error traceback\n"  
    "- Include all imports\n"  
    "- Use type hints\n"  
    "- For tensor shape mismatches: trace shapes step by step through the forward pass\n"  
    "- For import errors: check the dependent files for correct class/function names\n")
```

906

```

    "\n## Output Format\n"
    f"Write only the file: {crash_file}\n"
    "```\npython\n"
    "... \n"
    "```\n"
)

```

907

908 B.3.8 Validation of Intermediate Outputs

909 The pipeline produces several typed intermediate artefacts—gap lists, the architecture, the three
 910 independent plans, per-module specs, generated source, generated tests, the fix history—and validates
 911 each at a level matched to its downstream use rather than uniformly. The *architecture* is the only
 912 intermediate whose production is gated by an explicit machine check: the constraint validator
 913 described in §3.3.4 enforces orchestrator/module hygiene and is retried until it passes, and downstream
 914 stages refuse to run if it never does. *Generated source* is validated by `ast.parse` for syntactic well-
 915 formedness, with one self-correction call on parse failure; semantic correctness is not checked at
 916 generation time and surfaces only in the debug stage. *Per-module specs* are reviewed by the single
 917 review-pass LLM call in §3.3.4, whose output is re-injected into every code-generation prompt
 918 downstream—the pass is intended to surface interface inconsistencies before code is written, but
 919 it is not a machine check, and its precision is not measured within the pipeline. The remaining
 920 intermediates—gap lists, the three independent plans, generated tests—are not validated within the
 921 pipeline beyond the typed schema each is required to conform to; they are inputs to subsequent stages
 922 and are checked only insofar as a downstream consumer’s prompt fails to compose against them.
 923 End-to-end behavioural validation happens once, at the debug stage, against the entry point. This
 924 validation profile follows from the methodological principle in §3.1: each LLM call is given a typed,
 925 bounded context, and the *next* stage’s ability to consume that context is the implicit consistency check.
 926 Where a stage’s downstream consumer is itself an LLM call (rather than `ast.parse`, the constraint
 927 validator, or the subprocess), no further machine-level validation is performed; the responsibility
 928 for behavioural correctness is concentrated at the debug stage rather than distributed across every
 929 individual stage.

930 B.3.9 Failure Handling Across Stages

931 Each stage has bounded behaviour on pathological inputs rather than relying on the next stage
 932 to recover. **Architecture validator (§3.3.4)**: the architecture-design call is retried up to a fixed
 933 budget if the LLM returns an empty `files` array; orchestrator/module conflicts trigger a separate
 934 replan call (with its own smaller budget), and after that budget is exhausted the validator auto-fixes
 935 residual conflicts by stripping `implements` from orchestrator files, preferring forward progress with
 936 a corrected architecture over halting on a hygiene error the validator can resolve mechanically. **Gap
 937 stage (§3.3.3)**: if the gap-detection LLM call returns an empty list—interpreted as “*no gaps found
 938 in this paper*”—the pipeline proceeds with `supplementary_context` populated as empty fields,
 939 and planning operates as if the paper were fully self-contained. This is permissive by design, since
 940 no downstream prompt depends on `supplementary_context` being non-empty. **Code-generation
 941 self-correction (§3.3.5)**: a file whose self-correction also fails to parse is committed to disk regardless
 942 of the parse error; the `SyntaxError` then surfaces in the debug stage and is handled by the standard
 943 debug loop. **Debug subprocess without a parseable traceback (§3.3.7)**: the captured `stdout`
 944 and `stderr` are passed to the LLM with the crash file unset, and the model is expected to identify
 945 the offending file via the `search_codebase` tool before issuing a `submit_fix`. **Debug response
 946 without `submit_fix` or with a malformed tool call (§3.3.7)**: the iteration is treated as a no-op
 947 and the loop counter advances, bounding the cost of repeated failures at `max_iterations` rather
 948 than allowing the agent to spin without effect. **Fast-path `pip install` failure (§3.3.7)**: the failure
 949 message is appended to the captured output and the missing-module error re-enters the LLM-driven
 950 path on the next iteration, where the model can propose an alternative import or rewrite the offending
 951 file to remove the dependency.

952 B.4 Cross-Cutting Principles, Recapped

953 Section 3.4 audits the methodological principle stated at the start of Section 3 across the four
 954 places it applies, in increasing scope: from prompt-level (no-fabrication discipline), to file-level

(orchestrator/module separation), to context-level (closed-context discipline), to tool-level (multi-granularity debug primitives).

No-fabrication discipline as a layered invariant. The same constraint—*do not produce a value the paper does not state*—is enforced as a system-prompt-level instruction at every stage where a numerical, formulaic, or symbolic value would otherwise be invented: gap detection (the model must *flag* a missing value rather than guess it), goal-and-config extraction (the model must *transcribe* values from the paper rather than infer them), architecture design (files may not be *attributed* values not in the paper), sub-planning (formulas, weight-initialisation schemes, and named hyperparameters must be specified at the granularity the paper itself uses), and code generation (hyperparameters must be *read* from the configuration document rather than fabricated inline). The instruction is phrased differently at each stage because the constraint domain differs—the unit at risk of fabrication is a tag at gap detection, a row in a config at extraction, a file-level claim at architecture, an equation at sub-planning, and a Python literal at codegen. Distributing the rule across five stages rather than concentrating it at one chokepoint reflects a design choice rather than a measured property: a single chokepoint cannot enforce a constraint that applies to outputs at five different granularities. We do not claim any single layer is sufficient, and whether the layered invariant is empirically more reliable than a single chokepoint is a question for downstream evaluation rather than for the system description.

Orchestrator/module separation as a structural contract. The architecture validator described in §3.3.4 auto-tags entry points and `reproduce.sh` as orchestrators and strips any algorithmic responsibility from them. The same distinction propagates through every later stage: the sub-planner emits a different spec format for orchestrators (imports, sequence of operations, CLI parsing—no algorithms), and the code-generation agent assembles a different context for orchestrators (full source of every module rather than only declared dependencies). The combined effect is to prevent the common failure where the entry point silently re-implements model logic that already lives in a module, or where boilerplate scattered across multiple files double-implements the same equation.

Closed-context discipline as a fabrication-by-extrapolation filter. Two stages constrain the LLM by reducing what it is allowed to refer to rather than by enlarging what it is told. Sub-planning sees the gaps relevant to its module and not the global gap list. Code generation sees only the source of files it declared as dependencies (or, for orchestrators, the source of the modules it will glue together). Unit-test generation sees a whitelist of exactly which symbols are importable, AST-extracted from the generated source. `bam` a tighter expected scope: in our setting, where the goal is for the model to compose against a fixed set of declared symbols rather than to disambiguate among many plausible options, narrower prompts removed observed failure modes (cross-module symbol invention, test-time imports of non-existent helpers) during development. We do not claim this generalises beyond paper-to-code: in tasks where disambiguation depends on broader context, the same restrictions could harm performance.

Multi-granularity tool primitives as the debug-agent vocabulary. The debug agent’s `submit_fix` tool deliberately exposes three different granularities of change—line-range edits, full-file rewrites, and `files_to_create` that spawn fresh code-generation agents—rather than collapsing them into a single “rewrite this file” primitive. The choice of granularity is left to the model, which lets a localised bug (an off-by-one, a missing import) be fixed by an edit that does not regress the rest of the file, while a structural bug (a missing module, a wrong base class) can still be addressed by replacement or by re-invoking the planner-aware code-generation path. Pairing this with the dependency-cascade regeneration triggered after repeated failures (§3.3.7) gives the agent a smooth ladder from cheap, local repairs up to coarse, planner-conditioned regenerations.

1000 B.5 Memory Architecture

The system maintains two kinds of memory store: a single shared `GlobalMemory` and a per-agent `LocalMemory`. `GlobalMemory` is the only inter-agent communication channel: agents do not call each other directly, do not return values to a coordinator, and do not pass arguments to each other. Every agent reads its inputs from a fixed set of named keys in `GlobalMemory` and writes its outputs back to the same store. The keys are namespaced by stage of origin—`partitions` and `images` from partitioning; `figure_analysis` from the vision-language stage; `supplementary_context` from the gap stage; `architecture`, `goal_and_config`, `paper_analysis`, `experiment_plan`, `baseline_plan`, `dataset_plan`, `module_specs`, and `review` from planning; `generated_code` and `generated_tests` from code generation and unit-test generation—and the consumer of each

key is determined statically rather than dynamically dispatched. This decoupling is what makes individual stages re-runnable in isolation: invoking the debug agent against an existing run only requires that `generated_code` and `architecture` are populated, regardless of which other stages have been executed in the current process. On disk, `GlobalMemory` is one file per key, serialised as JSON for structured values (lists, dicts, the architecture, the generated source map) and as Markdown for free-text values (the goal-and-config document, the three plans, the review). The file-per-key layout means a partial state survives a process crash and can be reloaded by a fresh run that targets the same output directory; it also means that human inspection of any single key is a flat-file read rather than a parse of a monolithic blob.

Writes from the `asyncio` main loop have to be safe against concurrent puts from sub-agents that run under `asyncio.gather` (§3.2). The write path is split between the calling task and a background daemon thread. The calling task takes the value, decides whether the key is structured or textual, and serialises it to a string in the `asyncio` thread. It then spawns a daemon thread which acquires a per-store lock and writes the immutable string to disk. Because serialisation happens before the data crosses the thread boundary, the writer thread never observes a mutating dictionary, and because the writer holds the lock for the duration of the I/O, two concurrent writes to different keys cannot interleave their bytes. `LocalMemory` uses the same write pattern but is per-agent rather than shared, and contains three fields—`conversation_history`, `reasoning_trace`, and `intermediate_attempts`—populated through helper methods inherited from the base agent. `LocalMemory` is private by convention: outside the owning agent it is read in only two places—by the unit-test agent, which loads the codegen `LocalMemory` of each file to recover the spec it was written against, and by the debug agent, which can reconstruct the rationale for a particular generated file when diagnosing a failure. In every other case `LocalMemory` is write-only from the agent’s perspective and read-only from outside, which is what makes it safe to expose alongside the generated code at the end of a run without complicating the pipeline’s own dataflow.

B.6 Implementation Details

The pipeline is implemented in Python across the eight agent types described in §3.1 and their supporting modules. The default main-model backend is `MiniMax-M2.7-highspeed` accessed through an Anthropic-compatible endpoint, chosen for its 200k-token context window and its native tool-calling support, which the debug agent depends on. Four additional backends—Anthropic Claude, OpenAI GPT, DeepSeek, and Google Gemini—are wired into the same client and selected by a process-global `set_backend` call before `run_pipeline`, which makes per-backend comparisons a one-line change rather than a code edit. The vision-language stage is bound to Gemini 2.5-flash through its OpenAI-compatible endpoint and is independent of the main-model selection, so a comparison across main-model backends does not perturb the figure-analysis cost. Native tool-calling is implemented only for the Anthropic-style backends in our client; the OpenAI-style backends fall back to plain text completion in our setup, and the debug agent therefore requires an Anthropic-style backend to run. Generated code is executed as a Python subprocess in the host environment with the codebase directory as its working directory; missing third-party packages encountered during a debug run are installed on demand by `pip install` (§3.3.7) rather than baked into the host image ahead of time.

A batch runner drives the pipeline across multiple papers and backends in a single invocation: it switches the backend, resets the token-usage accumulator, runs the pipeline against each input directory, and writes one summary record per (backend × paper) pair. Slurm scripts in the repository wrap this runner for cluster execution and for the partial-pipeline configurations used in our experiments—most importantly `-skip-debug` for the four-stage variant and `-debug-iterations {20, 40}` for the two debug-budget variants. Within a single run, the pipeline is incrementally restartable: every stage that finds its output keys already populated in `GlobalMemory` will skip itself, so that a previously failed run can be resumed by pointing a new invocation at the same output directory and either re-running the failed stage in isolation or letting the cached upstream stages flow through unchanged. Combined with the file-per-key `GlobalMemory` layout described in §3.5 and the per-agent `LocalMemory` traces, this means every run leaves on disk a fully reconstructible record of what was generated, by which agent, at what cost, and against what context—the artefact set we use for the analyses in §4.

Table 7: Per-paper PaperBench Code-Dev scores in the Claude setting. All methods use Claude Sonnet 4.6 as the backend and are judged by o4-mini. Bold indicates the best score on each paper.

Paper	Paper2Code	DeepCode	Crafter
adaptive-pruning[55]	0.8604	0.4667	0.5889
all-in-one[12]	0.8667	0.7758	0.9403
bam[5]	0.8143	0.9456	0.9798
bbox[44]	0.7223	0.5855	0.8942
bridging-data-gaps[47]	0.4286	0.4821	0.8393
fre[11]	0.7251	0.6834	0.7099
ftrl[49]	0.3332	0.5728	0.7149
lbs[51]	0.8826	0.8695	0.8887
lca-on-the-line[40]	0.8178	0.7410	0.9124
mechanistic-understanding[23]	0.8426	0.9157	0.9398
pinn[6]	0.8593	0.8686	0.9506
rice[8]	0.3809	0.2118	0.8130
robust-clip[35]	0.0851	0.2613	0.8058
sample-specific-masks[4]	0.7978	0.7829	0.8326
sapg[42]	0.4888	0.4816	0.7774
self-composing-policies[26]	0.7712	0.6733	0.8315
self-expansion[46]	0.5986	0.4423	0.8051
semantic-self-consistency[21]	0.9626	0.9514	0.9522
sequential-neural-score-estimation[38]	0.9512	0.8418	0.9629
stay-on-topic-with-classifier-free-guidance[34]	0.8415	0.8605	0.8300
stochastic-interpolants[1]	0.6987	0.8370	0.7964
test-time-model-adaptation[28]	0.5766	0.8226	0.9103
what-will-my-model-forget[19]	0.7457	0.6682	0.5395

C Per-Paper Results on PaperBench Code-Dev

Tables 7 and 8 report the per-paper PaperBench Code-Dev scores for all methods. These tables provide the paper-level scores used to compute the win counts in Table 1. Bold entries indicate the best score on each paper within the same backend setting.

D Execution-Oriented Evaluation Details

This appendix provides additional details for the execution-oriented evaluation in Section 4.3. We describe the paper-selection criteria, the locked execution protocols, the fixture construction, the harness mechanics, the metric-extraction procedure, the repair-audit protocol, and the final per-cell execution records.

D.1 Paper Selection Criteria

We select three papers from the 23-paper PaperBench subset used in our Code-Dev evaluation. The selection is based on four criteria.

First, the empirical claim must be *contractable*: the paper’s central contribution can be expressed as a directional comparison between a method and a baseline, such as $\text{metric}_{\text{method}} \geq \text{metric}_{\text{baseline}}$. This allows the harness to distinguish between two cases: a repository that merely runs and emits numbers, and a repository that reproduces the qualitative direction of the paper’s main claim.

Second, the claim must be meaningful under fixture-scale execution. We require that the reduced experiment can run with at most one epoch and roughly 10^3 samples, so that smoke and pilot runs still test a real computational path rather than only training duration.

Third, the paper must have lightweight infrastructure requirements. We exclude papers that require external LLM APIs, specialized reinforcement-learning environments, or downloads of models larger than 1B parameters. This keeps all execution runs within a comparable single-GPU budget.

Fourth, the task must not introduce safety-refusal confounds. We exclude dual-use security papers involving adversarial attacks, jailbreaks, or red-teaming, because code-generation models may

Table 8: Per-paper PaperBench Code-Dev scores in the MiniMax setting. All methods use MiniMax M2.7 as the backend and are judged by gpt-4o-mini. Bold indicates the best score on each paper.

Paper	Paper2Code	DeepCode	Crafter
adaptive-pruning	0.3103	0.4406	0.5344
all-in-one	0.5499	0.4291	0.6008
bam	0.6681	0.3483	0.8132
bbox	0.3277	0.1795	0.5419
bridging-data-gaps	0.2768	0.4375	0.5315
fre	0.4109	0.5921	0.7579
fttl	0.3837	0.4588	0.6405
lbc	0.6163	0.5700	0.5861
lca-on-the-line	0.3794	0.2869	0.4614
mechanistic-understanding	0.3120	0.5046	0.8176
pinn	0.6330	0.7280	0.7847
rice	0.4121	0.3302	0.4908
robust-clip	0.3635	0.5930	0.6300
sample-specific-masks	0.6287	0.5186	0.4293
sapg	0.2363	0.2713	0.4063
self-composing-policies	0.6693	0.2431	0.5140
self-expansion	0.1835	0.3209	0.5190
semantic-self-consistency	0.5494	0.5510	0.6453
sequential-neural-score-estimation	0.7043	0.6649	0.6235
stay-on-topic-with-classifier-free-guidance	0.1180	0.3669	0.5719
stochastic-interpolants	0.6216	0.5341	0.7094
test-time-model-adaptation	0.1550	0.5076	0.6349
what-will-my-model-forget	0.1627	0.2852	0.2686

1088 refuse to emit such code. Otherwise, safety-policy behavior would become a confounder in the
1089 reproducibility evaluation.

1090 The resulting subset spans three different ML settings. LBCS tests data-efficient learning with bilevel
1091 optimization and lexicographic coreset selection. Self-Expansion (abbreviated as SEMA hereafter)
1092 tests continual learning with a frozen backbone, router, and dynamically expanding modular adapters.
1093 LCA-on-the-Line tests a semantic-distance evaluation mechanism through the paper’s controlled
1094 simulation phase.

1095 **Why LCA-on-the-Line uses Phase 1.** LCA-on-the-Line contains both a controlled simulation
1096 phase and a larger benchmark phase. Our protocol targets the controlled simulation phase, following
1097 Appendix C of the original paper. This phase is the mechanism check: it tests whether models using
1098 transferable causal features produce lower LCA distance than models using confounding features,
1099 even when their ID accuracy can be worse. The benchmark phase verifies this mechanism at scale
1100 across 75 pretrained vision models on ImageNet and OOD variants. That scale is outside our single-
1101 GPU smoke/pilot budget. Therefore, the Phase-1 protocol is the strongest scope-bounded executable
1102 test of the paper’s contribution direction.

1103 D.2 Locked Execution Protocols

1104 For each selected paper, we write one locked protocol before evaluating any generated repository. The
1105 protocol is paper-side and method-neutral. It specifies what work the repository must perform, but it
1106 does not prescribe a repository layout, command-line interface, configuration format, environment
1107 variables, or output path.

1108 Each protocol contains four components: a smoke intent, a pilot intent, a timeout budget, and a
1109 paper-level success predicate. The smoke run checks whether the repository can execute a minimal
1110 paper-relevant path and write at least one artifact. The pilot run executes a reduced but real version
1111 of the paper’s core comparison. We use a 600-second timeout for smoke runs and an 1800-second
1112 timeout for pilot runs.

1113 The success predicate is layered. M6 is a liveness milestone: the repository must emit the required
1114 finite scalar metrics in the protocol’s expected namespace. M7 is a contribution milestone: after M6

Table 9: Smoke and pilot intents in the locked execution protocols.

Paper	Smoke intent	Pilot intent
LBCS	Load a small MNIST-style dataset, initialize the paper-specified model, run one training epoch on a tiny coreset, and write at least one artifact.	Run a fixture-scale coreset-selection comparison on MNIST-S. The pilot must execute both LBCS and a uniform-random subset baseline at the same coreset size k , then emit their test accuracies side-by-side. Work may be reduced only through small k , a single seed, a single epoch, or a small batch size.
SEMA	Load the small multi-task vision fixture, initialize a small pre-trained backbone, run one forward pass over a single task, and write at least one artifact.	Run a fixture-scale continual-learning evaluation on <code>cifar_multitask</code> with 5 tasks and 2 classes per task. The pilot must execute the SEMA model through the task sequence, emit its final-task average accuracy A_N , and report the corresponding chance-level baseline. It may not use any flag that bypasses real adapter expansion.
LCA-on-the-Line	Load the hierarchical Gaussian-mixture fixture, train one simulation model on a single trial, compute ID top-1 error and LCA distance against the supplied hierarchy, and write at least one artifact.	Run the paper’s Phase-1 simulation on <code>gaussian_mixture</code> . The pilot must train both simulation models, f using transferable causal features and g using confounding features, on the same hierarchical Gaussian fixture. It must compute LCA distance for both models against the supplied taxonomy and emit the comparison side-by-side. It may not hard-code an alternative hierarchy that trivializes the LCA computation.

Table 10: Per-protocol liveness keys and contribution invariants. M6 checks finite, in-range emitted metrics; M7 checks the paper-specific contribution direction after M6 passes.

Paper	Liveness keys (M6)	Contribution invariant (M7)
LBCS	<code>acc_lbc</code> $\in (0, 100]$; <code>train_loss</code> finite and > 0	<code>acc_lbc</code> $>$ <code>acc_uniform</code> at the same k
SEMA	<code>acc_sema</code> $\in (0, 100]$; <code>final_task_loss</code> finite and > 0	<code>acc_sema</code> $>$ <code>acc_baseline</code> , where <code>acc_baseline</code> is the chance-level baseline for the evaluated task sequence
LCA-on-the-Line	<code>lca_method</code> finite and > 0 ; <code>id_top1_error</code> , <code>ood_top1_error</code> $\in [0, 1]$	<code>lca_method</code> $<$ <code>lca_baseline</code>

is satisfied, the repository must also satisfy the paper-specific directional invariant. M6 and M7 do not require matching the original paper’s reported absolute performance. They only require that the fixture-scale run produces a non-degenerate signal in the paper’s qualitative direction.

Table 9 lists the smoke and pilot intents. Table 10 lists the liveness keys and contribution invariants.

D.3 Fixtures

The harness uses deterministic fixture families shared across all evaluated methods. LBCS and SEMA use the `vision_classification` family, while LCA-on-the-Line uses the `synthetic_hierarchical` family.

For LBCS, `mnist_tiny` contains 1,000 MNIST images and is written both as a flat NPZ file and in `torchvision-compatible` MNIST format. For SEMA, `cifar_multitask` contains five two-class tasks sampled from CIFAR-100 and stored as an `ImageFolder`-style fixture. For LCA-on-the-Line, `gaussian_mixture` contains four-class hierarchical Gaussian data with ID and OOD splits, together with a two-level hierarchy file.

The harness supplies fixtures to the evaluated repository only through environment variables, including `MNIST_PATH`, `CIFAR_MULTITASK_PATH`, `GAUSSIAN_MIX_PATH`, and `HIERARCHY_FILE`. Since the fixture family is part of the locked paper protocol, all pipelines receive byte-identical inputs.

1131 D.4 Harness Mechanics

1132 The harness first probes each generated repository to summarize its available execution interface,
1133 including discovered entrypoints, command-line flags, environment variables, and configuration files.
1134 A fixed auxiliary planner then emits a JSON invocation plan for smoke and pilot runs. The plan is
1135 validated against this discovered interface: fabricated flags, unknown environment variables, and
1136 unregistered fixtures are rejected.

1137 After validation, the harness materializes the fixture and runs the selected repository entrypoint
1138 directly as a subprocess. It captures stdout and stderr without using a shell wrapper that could mask
1139 failures. The harness then scans newly written JSON artifacts, maps observed keys to canonical
1140 protocol metrics, and writes the evaluated `metrics.json` itself.

1141 This design prevents the planner or wrapper from fabricating success. The planner never edits
1142 repository code, the repository is invoked directly, and M6/M7 predicates are evaluated only over
1143 metrics extracted by the harness.

1144 The same locked protocol is used for the as-shipped evaluation and for every human-repair cycle.
1145 Thus, repair cycles do not change the target; they only change the generated repository being
1146 evaluated.

1147 D.5 Metric Extraction and Predicates

1148 Metric extraction is performed by the harness. The extractor scans JSON files under standard
1149 output directories, including `results/`, `outputs/`, `logs/`, `checkpoints/`, `figures/`, `runs/`, and
1150 `experiments/`. It flattens JSON leaf keys and first matches them against protocol-declared regex
1151 aliases.

1152 If a required key remains missing, a fixed LLM mapper may propose a dotted JSON path. The
1153 proposed path is resolved in process and accepted only if it points to a finite scalar. For each extracted
1154 metric, the harness records the source file, JSON path, and extraction method.

1155 The final predicates are simple directional checks. LBCS reaches M7 if `acc_lbc` > `acc_uniform`.
1156 SEMA reaches M7 if `acc_sema` > `acc_baseline`. LCA-on-the-Line reaches M7 if `lca_method` <
1157 `lca_baseline`.

1158 D.6 Repair Audit Protocol

1159 For cells that fail before repair, we run a bounded repair loop with a 20-cycle cap. Each cycle records
1160 the observed failure, root-cause category, patch description, files edited, and line delta. Automatic
1161 root-cause categories are computed from stdout and stderr using regex rules. These categories include
1162 syntax errors, missing imports, undefined symbols, function-signature mismatches, missing files or
1163 configs, dependency errors, CLI mismatches, tensor/device/dtype errors, timeouts, OOM failures,
1164 and invalid outputs.

1165 We additionally assign a manual severity tier to each repair patch. S1 denotes mechanical surface
1166 fixes, S2 denotes localized logic or configuration fixes, and S3 denotes architectural gaps, such as a
1167 missing paper phase, missing evaluation tail, or materially incorrect metric definition. Table 11 gives
1168 the full taxonomy.

1169 D.7 Final Emitted Metrics

1170 For transparency, we report the flat-key metrics emitted by each repaired repository on its final pilot
1171 run. These are the values extracted by the harness to evaluate M6 and M7. The numbers are read
1172 directly from the generated codebase’s output JSON; we do not manually post-process the values.

1173 **LBCS.** The contribution invariant is `acc_lbc` > `acc_uniform` at the same coreset size k .
1174 Table 12 reports the final emitted metrics for the LBCS pilot run.

1175 Only Crafter satisfies the LBCS contribution invariant at this fixture scale. The LBCS-over-uniform
1176 advantage reported in the original paper becomes unstable under our much smaller pilot setting,
1177 which uses a reduced MNIST-S fixture and small coreset sizes. Paper2Code and DeepCode therefore
1178 reach M6 but not M7.

Table 11: Repair-tier taxonomy used in the execution-oriented evaluation.

Tier	Category	LOC	Description / signature
S1	syntax error	1–3	<code>SyntaxError</code> , <code>IndentationError</code> , or <code>TabError</code> .
S1	missing import	1–3	<code>ModuleNotFoundError</code> , <code>ImportError</code> , or a missing dependency package.
S1	signature mismatch	1–3	Wrong or missing keyword argument; missing positional argument.
S1	missing file/config	1–3	<code>FileNotFoundError</code> ; missing path or configuration file.
S1	dependency setup	1–3	<code>pip</code> failure, version conflict, or <code>DistributionNotFound</code> .
S1	CLI mismatch	1–3	Unrecognized argument, invalid choice, or required flag missing.
S2	undefined symbol	5–30	<code>NameError</code> or <code>AttributeError</code> ; broken cross-module reference.
S2	tensor/device error	5–30	Mixed CPU/CUDA tensors, shape mismatch, or scalar-type mismatch.
S2	timeout/OOM	5–30	CUDA OOM, timeout, or killed process.
S2	no metrics	5–30	The run completes or partially completes, but does not emit the required finite scalar metrics.
S2	invalid output	5–30	Invalid shape, NaN propagation, or missing required scalar output.
S3	missing evaluation tail	30–100+	The main algorithm is implemented, but the evaluation or output stage is missing.
S3	missing paper phase	30–100+	The paper has multiple phases, but the codebase implements only a subset.
S3	incorrect metric	30–100+	A paper-cited metric function exists, but it computes the wrong quantity.

Table 12: Final emitted metrics for the LBCS fixture-scale pilot run. M7 requires `acc_lbc` > `acc_uniform` at the same coreset size k .

Pipeline	k	<code>acc_lbc</code> (%)	<code>acc_uniform</code> (%)	<code>train_loss</code>	Verdict
Crafter	50	14.71	13.00	1.000	M7 ($\Delta = +1.71$)
Paper2Code	197	11.35	11.36	0.500	M6
DeepCode	50	9.74	9.82	2.296	M6

1179 **SEMA.** The contribution invariant is `acc_sema` > `acc_baseline`, where `acc_baseline` is the
1180 chance-level accuracy for the evaluated task sequence. Table 13 reports the final emitted metrics for
1181 the SEMA pilot run.

1182 All three pipelines satisfy the SEMA contribution invariant after repair. DeepCode’s final
1183 `acc_sema`=28.30 is emitted from codebase training rather than from a hand-written harness value. It
1184 appears only after 12 repair cycles that resolved multiple cross-module naming inconsistencies in the
1185 generated repository.

1186 **LCA-on-the-Line.** The contribution invariant is `lca_method` < `lca_baseline`, where f uses
1187 transferable causal features and g uses confounding features. Table 14 reports the final emitted
1188 metrics for the Phase-1 simulation.

Table 13: Final emitted metrics for the SEMA fixture-scale pilot run. M7 requires `acc_sema` to exceed the chance-level baseline for the evaluated task sequence.

Pipeline	num_tasks	acc_sema (%)	acc_baseline (%)	final_task_loss
Crafter	2	2.69	1.84	1.000
Paper2Code	5	25.40	10.00	0.746
DeepCode	5	28.30	2.00	1.176

Table 14: Final emitted metrics for the LCA-on-the-Line Phase-1 simulation. M7 requires `lca_method < lca_baseline`.

Pipeline	lca_method (<i>f</i>)	lca_baseline (<i>g</i>)	id_top1_error	ood_top1_error
Crafter	1.00	2.00	0.159	0.159
Paper2Code	1.00	2.00	0.159	0.159
DeepCode	1.00	2.00	0.161	0.161

1189 All three pipelines reproduce the qualitative Phase-1 prediction: $LCA(f) < LCA(g)$. The exact
 1190 values 1.00 and 2.00 reflect the controlled four-class hierarchy, where LCA distances are integer-
 1191 valued. The *f*-model’s errors mostly remain within sibling classes, while the *g*-model’s errors cross
 1192 to non-sibling classes.

1193 D.8 Full Repair Records

1194 Table 15 reports the full per-cell execution-evaluation record. Pre-M is the milestone reached on the
 1195 first invocation. #cyc is the total number of repair cycles run, capped at 20. M7_cyc is the cycle
 1196 index at which M7 was first reached; NR means that M7 was not reached. $|\Delta LOC|$ is additions plus
 1197 deletions in `working/` relative to the generation-time `baseline/`. S1, S2, and S3 count the distinct
 1198 repair patches applied at each severity tier.

Table 15: Full per-cell execution-evaluation record.

Paper	Method	Pre-M	#cyc	M7 cyc	Final M	$ \Delta LOC $	S1	S2	S3
LBCS	Crafter	M5	3	2	M7	43	1	0	0
LBCS	Paper2Code	M1	13	NR	M6	151	4	0	1
LBCS	DeepCode	M1	13	NR	M6	165	3	1	1
SEMA	Crafter	M5	7	6	M7	137	1	2	0
SEMA	Paper2Code	M5	13	11	M7	116	3	2	1
SEMA	DeepCode	M1	12	12	M7	185	7	5	1
LCA	Crafter	M5	5	4	M7	52	0	0	0
LCA	Paper2Code	M5	8	7	M7	55	2	0	1
LCA	DeepCode	M1	11	8	M7	71	2	0	1

1199 NeurIPS Paper Checklist

1200 1. Claims

1201 Question: Do the main claims made in the abstract and introduction accurately reflect the
1202 paper’s contributions and scope?

1203 Answer: [Yes]

1204 Justification: The abstract and introduction accurately state the paper’s main contributions:
1205 an agentic code-generation pipeline for deep learning reproducibility and an execution-
1206 readiness evaluation protocol. The claims are scoped to the PaperBench setting and the
1207 experiments reported in the paper.

1208 Guidelines:

- 1209 • The answer [N/A] means that the abstract and introduction do not include the claims
1210 made in the paper.
- 1211 • The abstract and/or introduction should clearly state the claims made, including the
1212 contributions made in the paper and important assumptions and limitations. A [No] or
1213 [N/A] answer to this question will not be perceived well by the reviewers.
- 1214 • The claims made should match theoretical and experimental results, and reflect how
1215 much the results can be expected to generalize to other settings.
- 1216 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1217 are not attained by the paper.

1218 2. Limitations

1219 Question: Does the paper discuss the limitations of the work performed by the authors?

1220 Answer: [Yes]

1221 Justification: The paper discusses limitations in the Future Work section, including the
1222 remaining need for human bug fixing, the focus on qualitative rather than exact quantitative
1223 reproduction, and the challenge of runtime environment configuration. It also outlines these
1224 issues as directions for future work toward more automated reproducible ML research.

1225 Guidelines:

- 1226 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1227 that the paper has limitations, but those are not discussed in the paper.
- 1228 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1229 • The paper should point out any strong assumptions and how robust the results are to
1230 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1231 model well-specification, asymptotic approximations only holding locally). The authors
1232 should reflect on how these assumptions might be violated in practice and what the
1233 implications would be.
- 1234 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1235 only tested on a few datasets or with a few runs. In general, empirical results often
1236 depend on implicit assumptions, which should be articulated.
- 1237 • The authors should reflect on the factors that influence the performance of the approach.
1238 For example, a facial recognition algorithm may perform poorly when image resolution
1239 is low or images are taken in low lighting. Or a speech-to-text system might not be
1240 used reliably to provide closed captions for online lectures because it fails to handle
1241 technical jargon.
- 1242 • The authors should discuss the computational efficiency of the proposed algorithms
1243 and how they scale with dataset size.
- 1244 • If applicable, the authors should discuss possible limitations of their approach to
1245 address problems of privacy and fairness.
- 1246 • While the authors might fear that complete honesty about limitations might be used by
1247 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
1248 limitations that aren’t acknowledged in the paper. The authors should use their best
1249 judgment and recognize that individual actions in favor of transparency play an impor-
1250 tant role in developing norms that preserve the integrity of the community. Reviewers
1251 will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper does not present new theoretical results, theorems, or formal proofs. Its contributions are methodological and empirical.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper describes the benchmark, evaluated papers, baselines, backend models, judges, scoring protocol, execution-readiness harness, and ablation settings used for the main experiments. The appendix further provides paper-specific fixture protocols, metric keys, and contribution predicates needed to reproduce the evaluation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

1305 (d) We recognize that reproducibility may be tricky in some cases, in which case
1306 authors are welcome to describe the particular way they provide for reproducibility.
1307 In the case of closed-source models, it may be that access to the model is limited in
1308 some way (e.g., to registered users), but it should be possible for other researchers
1309 to have some path to reproducing or verifying the results.

1310 5. Open access to data and code

1311 Question: Does the paper provide open access to the data and code, with sufficient instruc-
1312 tions to faithfully reproduce the main experimental results, as described in supplemental
1313 material?

1314 Answer: [Yes]

1315 Justification: We submit the source code for the proposed pipeline and evaluation scripts
1316 with the supplemental material.

1317 Guidelines:

- 1318 • The answer [N/A] means that paper does not include experiments requiring code.
- 1319 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/
1320 public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 1321 • While we encourage the release of code and data, we understand that this might not
1322 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1323 including code, unless this is central to the contribution (e.g., for a new open-source
1324 benchmark).
- 1325 • The instructions should contain the exact command and environment needed to run to
1326 reproduce the results. See the NeurIPS code and data submission guidelines ([https:
1327 //neurips.cc/public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 1328 • The authors should provide instructions on data access and preparation, including how
1329 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1330 • The authors should provide scripts to reproduce all experimental results for the new
1331 proposed method and baselines. If only a subset of experiments are reproducible, they
1332 should state which ones are omitted from the script and why.
- 1333 • At submission time, to preserve anonymity, the authors should release anonymized
1334 versions (if applicable).
- 1335 • Providing as much information as possible in supplemental material (appended to the
1336 paper) is recommended, but including URLs to data and code is permitted.

1337 6. Experimental setting/details

1338 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1339 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1340 Answer: [Yes]

1341 Justification: The paper specifies the key experimental settings, including the benchmark,
1342 evaluated papers, baselines, backend models, judges, scoring protocol, ablation variants, and
1343 execution-readiness fixture design. Details such as fixture-scale inputs, metric keys, time
1344 limits, and contribution predicates are provided in the appendix.

1345 Guidelines:

- 1346 • The answer [N/A] means that the paper does not include experiments.
- 1347 • The experimental setting should be presented in the core of the paper to a level of detail
1348 that is necessary to appreciate the results and make sense of them.
- 1349 • The full details can be provided either with the code, in appendix, or as supplemental
1350 material.

1351 7. Experiment statistical significance

1352 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1353 information about the statistical significance of the experiments?

1354 Answer: [No]

Justification: The paper does not report formal error bars or statistical significance tests because the main results are based on benchmark scores and execution outcomes rather than repeated stochastic trials. We instead report results across different backend/judge settings and ablation variants to assess the robustness of the observed trends.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper reports the compute setup used for the main experiments, including GH200 GPU nodes for execution and repair evaluation. It also specifies the time limits for fixture-scale runs, including the smoke and pilot execution budgets.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conforms to the NeurIPS Code of Ethics: it does not involve human subjects, private or sensitive data, high-risk model release, or deployment in safety-critical settings. The paper focuses on improving ML reproducibility and reports the evaluation setup and limitations transparently.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [No]

Justification: The paper does not include a dedicated broader impact discussion. The work is primarily methodological and focuses on improving ML reproducibility, with no direct deployment or human-subject setting.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models, image generators, scraped datasets, or sensitive data. The work focuses on an evaluation and code-generation pipeline for ML reproducibility, so responsible-release safeguards for high-risk assets are not applicable.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper credits the existing benchmarks, baseline systems, datasets, and model APIs used in the experiments through citations and descriptions in the experimental setup. We respect the applicable access conditions and do not repackage or redistribute third-party assets beyond their permitted use.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release new datasets, models, or code assets at submission time. The proposed pipeline and evaluation protocol are described in the paper, but no standalone new asset is distributed with the submission.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing, user studies, or research with human subjects. Therefore, participant instructions, screenshots, and compensation details are not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

1514 • Including this information in the supplemental material is fine, but if the main contribu-
 1515 tion of the paper involves human subjects, then as much detail as possible should be
 1516 included in the main paper.

1517 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
 1518 or other labor should be paid at least the minimum wage in the country of the data
 1519 collector.

1520 **15. Institutional review board (IRB) approvals or equivalent for research with human**
 1521 **subjects**

1522 Question: Does the paper describe potential risks incurred by study participants, whether
 1523 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
 1524 approvals (or an equivalent approval/review based on the requirements of your country or
 1525 institution) were obtained?

1526 Answer: [N/A]

1527 Justification: The paper does not involve crowdsourcing, user studies, or research with
 1528 human subjects. Therefore, IRB approval or equivalent review is not applicable.

1529 Guidelines:

1530 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
 1531 with human subjects.

1532 • Depending on the country in which research is conducted, IRB approval (or equivalent)
 1533 may be required for any human subjects research. If you obtained IRB approval, you
 1534 should clearly state this in the paper.

1535 • We recognize that the procedures for this may vary significantly between institutions
 1536 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
 1537 guidelines for their institution.

1538 • For initial submissions, do not include any information that would break anonymity (if
 1539 applicable), such as the institution conducting the review.

1540 **16. Declaration of LLM usage**

1541 Question: Does the paper describe the usage of LLMs if it is an important, original, or
 1542 non-standard component of the core methods in this research? Note that if the LLM is used
 1543 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
 1544 scientific rigor, or originality of the research, declaration is not required.

1545 Answer: [Yes]

1546 Justification: LLMs are a core component of the proposed method and are used for speci-
 1547 fication construction, gap filling, code generation, and execution-aware repair. The paper
 1548 describes the LLM backends and judge models used in the main experiments.

1549 Guidelines:

1550 • The answer [N/A] means that the core method development in this research does not
 1551 involve LLMs as any important, original, or non-standard components.

1552 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
 1553 be described.